

# Riak Tutorial

Sean Cribbs & Ian Plosker  
Basho Technologies

# What is Riak?



- Key-Value store (plus extras)

- Key-Value store (plus extras)
- Distributed, horizontally scalable

- Key-Value store (plus extras)
- Distributed, horizontally scalable
- Fault-tolerant

- Key-Value store (plus extras)
- Distributed, horizontally scalable
- Fault-tolerant
- Highly-available

- Key-Value store (plus extras)
- Distributed, horizontally scalable
- Fault-tolerant
- Highly-available
- Built for the Web



- Key-Value store (plus extras)
- Distributed, horizontally scalable
- Fault-tolerant
- Highly-available
- Built for the Web
- Inspired by Amazon's Dynamo

# Key-Value

# Key-Value

- Simple operations - get, put, delete

# Key-Value

- Simple operations - get, put, delete
- Value is mostly opaque (some metadata)

# Key-Value

- Simple operations - get, put, delete
- Value is mostly opaque (some metadata)
- Extras

# Key-Value

- Simple operations - get, put, delete
- Value is mostly opaque (some metadata)
- Extras
  - MapReduce

# Key-Value

- Simple operations - get, put, delete
- Value is mostly opaque (some metadata)
- Extras
  - MapReduce
  - Links

# Key-Value

- Simple operations - get, put, delete
- Value is mostly opaque (some metadata)
- Extras
  - MapReduce
  - Links
  - Full-text search (new, optional)



# Key-Value

- Simple operations - get, put, delete
- Value is mostly opaque (some metadata)
- Extras
  - MapReduce
  - Links
  - Full-text search (new, optional)
  - Secondary Indexes

# Distributed & Horizontally Scalable

# Distributed & Horizontally Scalable

- Default configuration is in a cluster

# Distributed & Horizontally Scalable

- Default configuration is in a cluster
- Load and data are spread evenly

# Distributed & Horizontally Scalable

- Default configuration is in a cluster
- Load and data are spread evenly
- Add more nodes to get more X

# Fault-Tolerant

# Fault-Tolerant

- All nodes participate equally - no SPOF

# Fault-Tolerant

- All nodes participate equally - no SPOF
- All data is replicated



# Fault-Tolerant

- All nodes participate equally - no SPOF
- All data is replicated
- Cluster transparently survives...

# Fault-Tolerant

- All nodes participate equally - no SPOF
- All data is replicated
- Cluster transparently survives...
  - node failure

# Fault-Tolerant

- All nodes participate equally - no SPOF
- All data is replicated
- Cluster transparently survives...
  - node failure
  - network partitions

# Fault-Tolerant

- All nodes participate equally - no SPOF
- All data is replicated
- Cluster transparently survives...
  - node failure
  - network partitions
- Built on Erlang/OTP (designed for FT)

# Highly-Available

# Highly-Available

- Any node can serve client requests

# Highly-Available

- Any node can serve client requests
- Fallbacks are used when nodes are down

# Highly-Available

- Any node can serve client requests
- Fallbacks are used when nodes are down
- Always accepts read and write requests



# Highly-Available

- Any node can serve client requests
- Fallbacks are used when nodes are down
- Always accepts read and write requests
- Per-request quorums

# Built for the Web

# Built for the Web

- HTTP is default (but not only) interface

# Built for the Web

- HTTP is default (but not only) interface
- Does HTTP/REST well (see Webmachine)

# Built for the Web

- HTTP is default (but not only) interface
- Does HTTP/REST well (see Webmachine)
- Plays well with HTTP infrastructure - reverse proxy caches, load balancers, web servers

# Built for the Web

- HTTP is default (but not only) interface
- Does HTTP/REST well (see Webmachine)
- Plays well with HTTP infrastructure - reverse proxy caches, load balancers, web servers
- Suitable to many web applications

# Inspired by Amazon's Dynamo

# Inspired by Amazon's Dynamo

- Masterless, peer-coordinated replication



# Inspired by Amazon's Dynamo

- Masterless, peer-coordinated replication
- Consistent hashing

# Inspired by Amazon's Dynamo

- Masterless, peer-coordinated replication
- Consistent hashing
- Eventually consistent

# Inspired by Amazon's Dynamo

- Masterless, peer-coordinated replication
- Consistent hashing
- Eventually consistent
- Quorum reads and writes

# Inspired by Amazon's Dynamo

- Masterless, peer-coordinated replication
- Consistent hashing
- Eventually consistent
- Quorum reads and writes
- Anti-entropy: read repair, hinted handoff

# Lab: Riak Basics

# Installing Riak

# Installing Riak

- Download from our local server

# Installing Riak

- Download from our local server
- Download a package:  
<http://downloads.basho.com/riak/CURRENT>



# Installing Riak

- Download from our local server
- Download a package:  
<http://downloads.basho.com/riak/CURRENT>
- Clone & build (need git and Erlang) :  

```
$ git clone git://github.com/basho/riak.git  
$ make all devrel
```

# You need `curl`

`apt-get install curl`

`yum install curl`

# Start the Cluster

# Start the Cluster

```
$ surge-start.sh
```

# Start the Cluster

```
$ surge-start.sh  
dev1/bin/riak start  
dev2/bin/riak start  
dev3/bin/riak start  
dev4/bin/riak start
```

# Start the Cluster

```
$ surge-start.sh
dev1/bin/riak start
dev2/bin/riak start
dev3/bin/riak start
dev4/bin/riak start
dev2/bin/riak-admin join dev1
Sent join request to dev1
dev3/bin/riak-admin join dev1
Sent join request to dev1
dev4/bin/riak-admin join dev1
Sent join request to dev1
```

# Check its status

# Check its status

```
$ dev1/bin/riak-admin member_status
```

```
$ dev1/bin/riak-admin ring_status
```

```
$ dev1/bin/riak-admin status
```



# PUT

# PUT

```
$ surge-put.sh
```

# PUT

```
$ surge-put.sh  
Enter a key: your-key
```

# PUT

```
$ surge-put.sh  
Enter a key: your-key  
Enter a value: your-value
```

# PUT

```
$ surge-put.sh  
Enter a key: your-key  
Enter a value: your-value  
Enter a write quorum value:
```

# PUT

```
$ surge-put.sh  
Enter a key: your-key  
Enter a value: your-value  
Enter a write quorum value:
```

```
* About to connect() to 127.0.0.1 port 8091 (#0)  
*   Trying 127.0.0.1... connected  
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)  
> PUT /riak/surge/your-key?returnbody=true&w=quorum HTTP/1.1  
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r  
zlib/1.2.5  
> Host: 127.0.0.1:8091  
> Accept: */*  
> Content-Type: text/plain  
> Content-Length: 10  
>
```

# PUT

```
$ surge-put.sh
Enter a key: your-key
Enter a value: your-value
Enter a write quorum value:

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/your-key?returnbody=true&w=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
> Content-Type: text/plain
> Content-Length: 10
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1qb3LYEpkzGNleL/k23G+LAA=
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Date: Tue, 27 Sep 2011 19:21:08 GMT
< Content-Type: text/plain
< Content-Length: 10
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
your-value
```

# GET



# GET

```
$ surge-get.sh
```

# GET

```
$ surge-get.sh  
Enter a key: your-key
```

# GET

```
$ surge-get.sh  
Enter a key: your-key  
Enter a read quorum value (default: quorum):
```

# GET

```
$ surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
```

# GET


```
$ surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1qb3LYEpkzGNleL/k23G+LAA=
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Last-Modified: Tue, 27 Sep 2011 19:01:45 GMT
< ETag: "51h3q7RjTNaHWYp04P0MJj"
< Date: Tue, 27 Sep 2011 19:31:01 GMT
< Content-Type: text/plain
< Content-Length: 10
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
your-value
```

# GET

```
$ surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1qb3LYEpkzGNleL/k23G+LAA=
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Last-Modified: Tue, 27 Sep 2011 19:01:45 GMT
< ETag: "51h3q7RjTNaHWYp04P0MJj"
< Date: Tue, 27 Sep 2011 19:31:01 GMT
< Content-Type: text/plain
< Content-Length: 10
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
your-value
```




vector clock

# GET

```
$ surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1qb3LYEpkzGNleL/k23G+LAA=
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Last-Modified: Tue, 27 Sep 2011 19:01:45 GMT
< ETag: "51h3q7RjTNaHWYp04P0MJj"
< Date: Tue, 27 Sep 2011 19:31:01 GMT
< Content-Type: text/plain
< Content-Length: 10
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
your-value
```



caching headers

# GET

```
$ surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1qb3LYEpkzGNleL/k23G+LAA=
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Last-Modified: Tue, 27 Sep 2011 19:01:45 GMT
< ETag: "51h3q7RjTNaHWYp04P0MJj"
< Date: Tue, 27 Sep 2011 19:31:01 GMT
< Content-Type: text/plain
< Content-Length: 10
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
your-value
```

content-type



# DELETE

# DELETE

```
$ surge-delete.sh
```

# DELETE

```
$ surge-delete.sh  
Type a key: your-key
```

# DELETE

```
$ surge-delete.sh
Type a key: your-key

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> DELETE /riak/surge/your-key HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
```

# DELETE

```
$ surge-delete.sh
Type a key: your-key

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> DELETE /riak/surge/your-key HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
< HTTP/1.1 204 No Content
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Date: Tue, 27 Sep 2011 19:33:59 GMT
< Content-Type: text/plain
< Content-Length: 0
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
```

# Kill two nodes

# Kill two nodes

```
$ dev4/bin/riak stop
```

# Kill two nodes

```
$ dev4/bin/riak stop  
ok
```



# Kill two nodes

```
$ dev4/bin/riak stop  
ok
```

```
$ dev3/bin/riak stop
```

# Kill two nodes

```
$ dev4/bin/riak stop  
ok
```

```
$ dev3/bin/riak stop  
ok
```

# Try GETting your key

# Try GETting your key

```
$ . ~/Desktop/surge/surge-get.sh
```

# Try GETting your key

```
$ . ~/Desktop/surge/surge-get.sh  
Enter a key: your-key
```

# Try GETting your key

```
$ . ~/Desktop/surge/surge-get.sh  
Enter a key: your-key  
Enter a read quorum value (default: quorum):
```

# Try GETting your key

```
$ . ~/Desktop/surge/surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
```

# Try GETting your key

```
$ . ~/Desktop/surge/surge-get.sh
Enter a key: your-key
Enter a read quorum value (default: quorum):

* About to connect() to 127.0.0.1 port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to 127.0.0.1 (127.0.0.1) port 8091 (#0)
> GET /riak/surge/your-key?r=quorum HTTP/1.1
> User-Agent: curl/7.21.4 (universal-apple-darwin11.0) libcurl/7.21.4 OpenSSL/0.9.8r
zlib/1.2.5
> Host: 127.0.0.1:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1qb3LYEpkzGNleL/k23G+LAA=
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Last-Modified: Tue, 27 Sep 2011 19:50:07 GMT
< ETag: "51h3q7RjTNaHWYp04P0MJj"
< Date: Tue, 27 Sep 2011 19:51:10 GMT
< Content-Type: text/plain
< Content-Length: 10
<
* Connection #0 to host 127.0.0.1 left intact
* Closing connection #0
your-value
```




# Riak Architecture

# Key-Value Data Model


# Key-Value Data Model

- All data (objects) are referenced by keys




# Key-Value Data Model

- All data (objects) are referenced by keys
- Keys are grouped into buckets



# Key-Value Data Model

- All data (objects) are referenced by keys
- Keys are grouped into buckets
- Simple operations: get, put, delete



# Key-Value Data Model

- All data (objects) are referenced by keys
- Keys are grouped into buckets
- Simple operations: get, put, delete
- Object is composed of metadata and value

**people/beth**

vclock: ...

content-type: text/html

last-mod: 20101108T...


links: ...

<html><head>....

# Consistent Hashing & The Ring

# Consistent Hashing & The Ring


- 160-bit integer keyspace





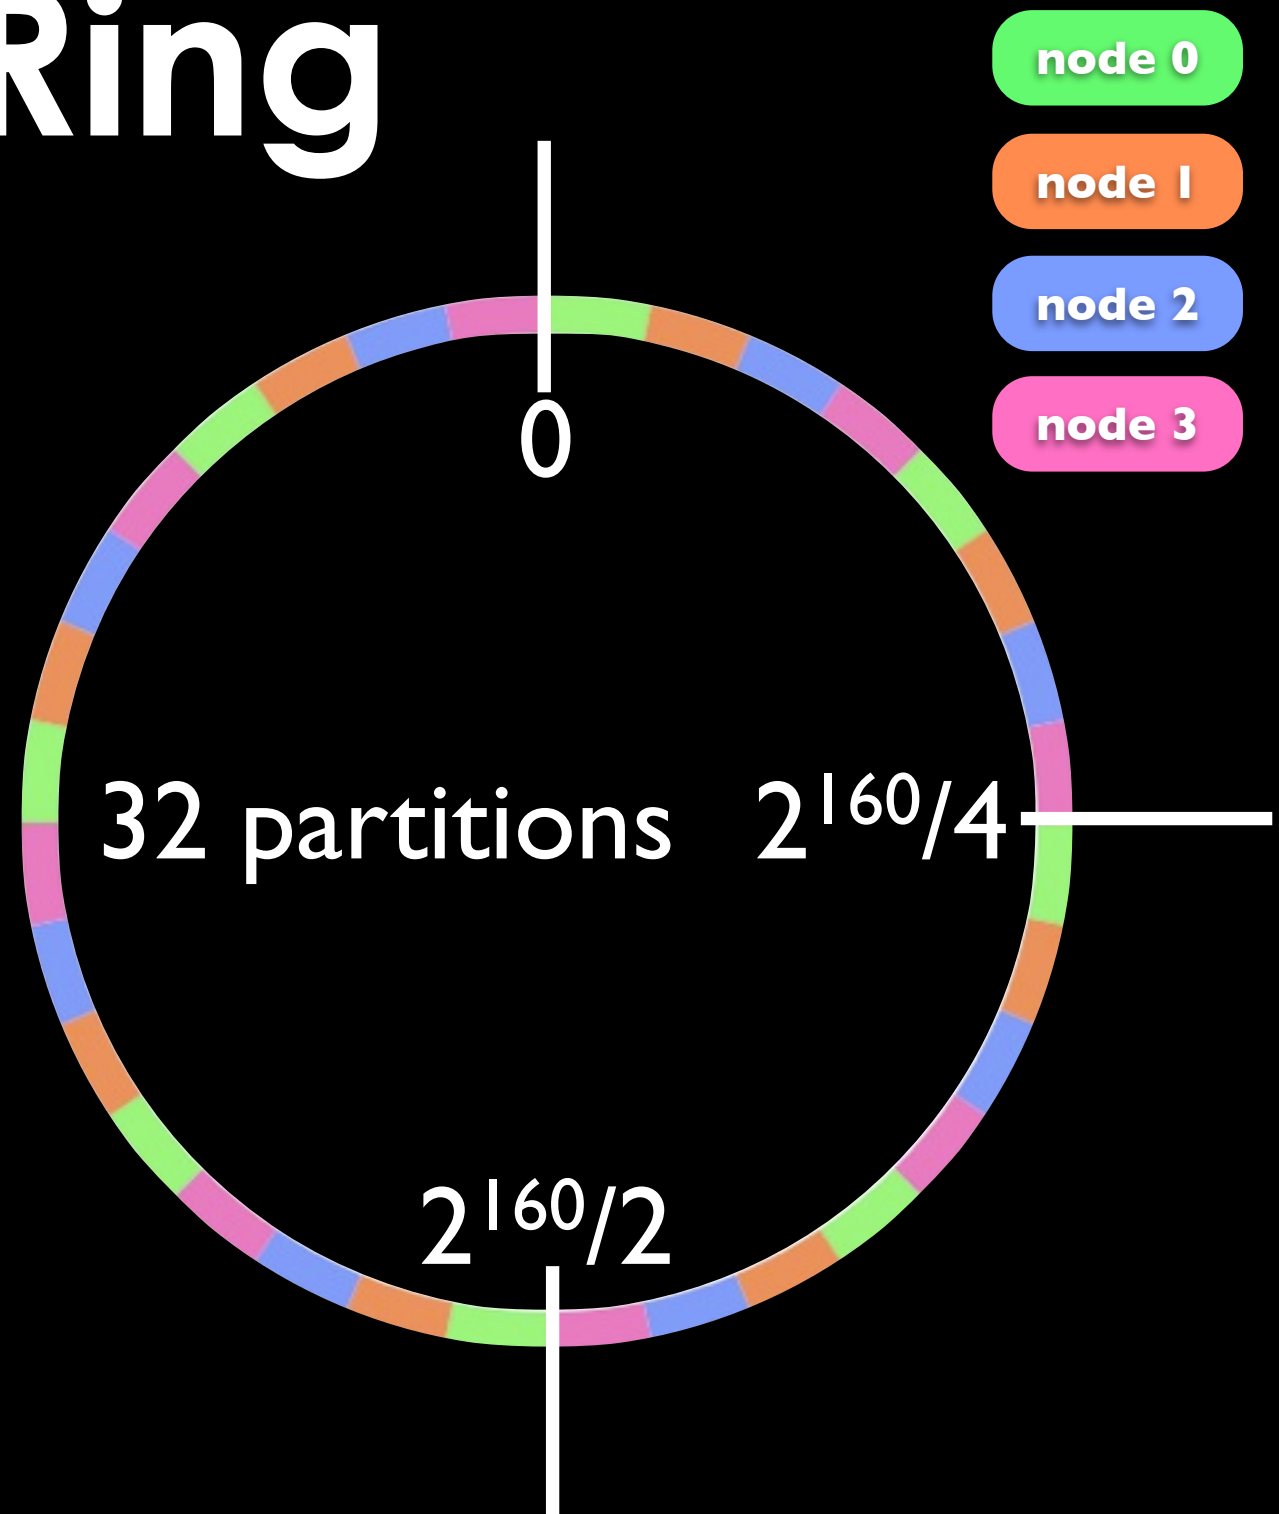
# Consistent Hashing & The Ring

- 160-bit integer keyspace
- Divided into fixed number of evenly-sized partitions



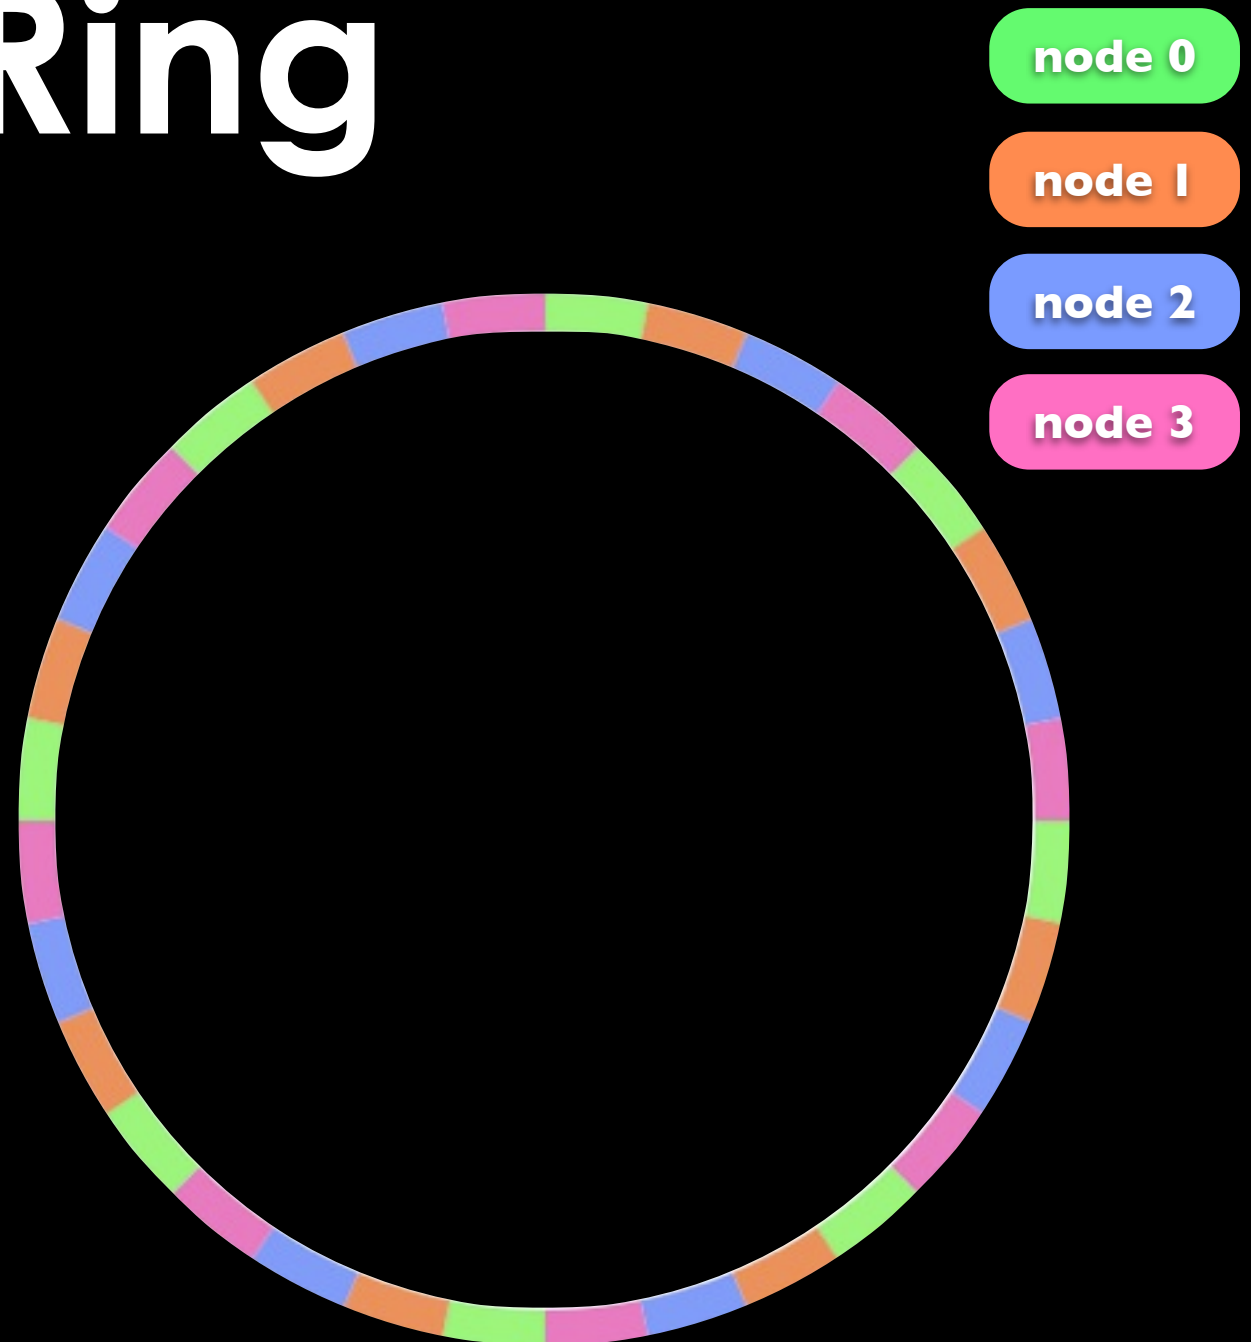
# Consistent Hashing & The Ring

- 160-bit integer keyspace
- Divided into fixed number of evenly-sized partitions
- Partitions are claimed by nodes in the cluster



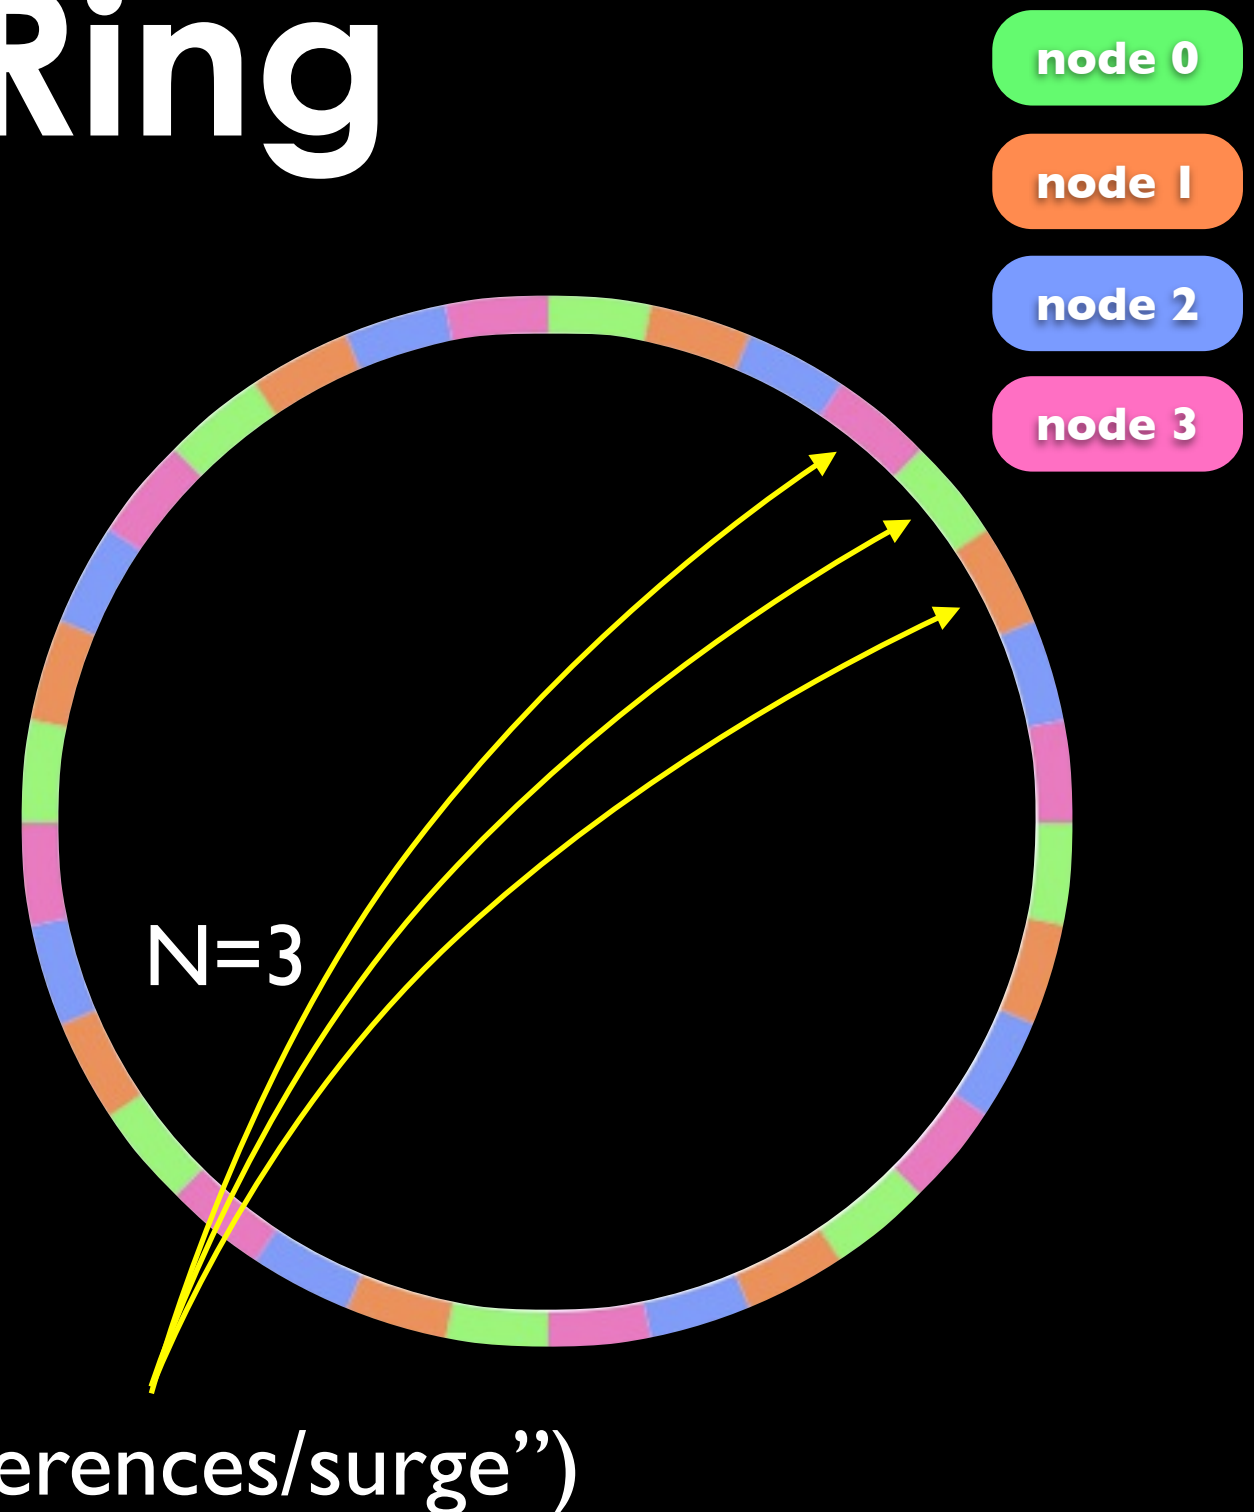
# Consistent Hashing & The Ring

- 160-bit integer keyspace
- Divided into fixed number of evenly-sized partitions
- Partitions are claimed by nodes in the cluster
- Replicas go to the N partitions following the key



# Consistent Hashing & The Ring

- 160-bit integer keyspace
- Divided into fixed number of evenly-sized partitions
- Partitions are claimed by nodes in the cluster
- Replicas go to the N partitions following the key



# Request Quorums

# Request Quorums

- Every request contacts all replicas of key

# Request Quorums

- Every request contacts all replicas of key
- N - number of replicas (default 3)

# Request Quorums

- Every request contacts all replicas of key
- $N$  - number of replicas (default 3)
- $R$  - read quorum



# Request Quorums

- Every request contacts all replicas of key
- $N$  - number of replicas (default 3)
- $R$  - read quorum
- $W$  - write quorum

# Vector Clocks

# Vector Clocks

- Every node has an ID (*changed in 1.0*)

# Vector Clocks

- Every node has an ID (*changed in 1.0*)
- Send last-seen vector clock in every “put” or “delete” request

# Vector Clocks

- Every node has an ID (*changed in 1.0*)
- Send last-seen vector clock in every “put” or “delete” request
- Riak tracks history of updates

# Vector Clocks

- Every node has an ID (*changed in 1.0*)
- Send last-seen vector clock in every “put” or “delete” request
- Riak tracks history of updates
  - Auto-resolves stale versions

# Vector Clocks

- Every node has an ID *(changed in 1.0)*
- Send last-seen vector clock in every “put” or “delete” request
- Riak tracks history of updates
  - Auto-resolves stale versions
  - Lets you decide conflicts


# Vector Clocks

<http://blog.basho.com/2010/01/29/why-vector-clocks-are-easy/>

- Every node has an ID (*changed in 1.0*)
- Send last-seen vector clock in every “put” or “delete” request
- Riak tracks history of updates
  - Auto-resolves stale versions
  - Lets you decide conflicts




# Hinted Handoff




# Hinted Handoff

- Node fails



# Hinted Handoff


- Node fails
- Requests go to fallback



hash("conferences/surge")

# Hinted Handoff


- Node fails
- Requests go to fallback
- Node comes back



hash("conferences/surge")

# Hinted Handoff


- Node fails
- Requests go to fallback
- Node comes back
- “Handoff” - data returns to recovered node



hash(“conferences/surge”)

# Hinted Handoff

- Node fails
- Requests go to fallback
- Node comes back
- “Handoff” - data returns to recovered node
- Normal operations resume



hash("conferences/surge")

# Anatomy of a Request

`get("conferences/surge")`

# Anatomy of a Request

get("conferences/surge")

client

Riak





# Anatomy of a Request

get("conferences/surge")

client  
Riak



Get Handler (FSM)

# Anatomy of a Request

get("conferences/surge")

client

Riak

Get Handler (FSM)

hash("conferences/oredev")  
== 10, 11, 12

# Anatomy of a Request

get("conferences/surge")

client


Riak

Get Handler (FSM)

hash("conferences/oredev")  
== 10, 11, 12

Coordinating node

Cluster



The Ring

# Anatomy of a Request

get("conferences/surge")

client


Riak

Get Handler (FSM)

get("conferences/oredev")

Coordinating node

Cluster



The Ring

# Anatomy of a Request

get("conferences/surge")

client


Riak

R=2

Get Handler (FSM)

Coordinating node

Cluster



The Ring

# Anatomy of a Request

get("conferences/surge")

client

Riak


R=2

v1

Get Handler (FSM)

Coordinating node

Cluster



The Ring

# Anatomy of a Request

get("conferences/surge")

client

Riak

R=2

v1

v2

Get Handler (FSM)

# Anatomy of a Request

get("conferences/surge")

client  
Riak

v2

R=2

v2

Get Handler (FSM)

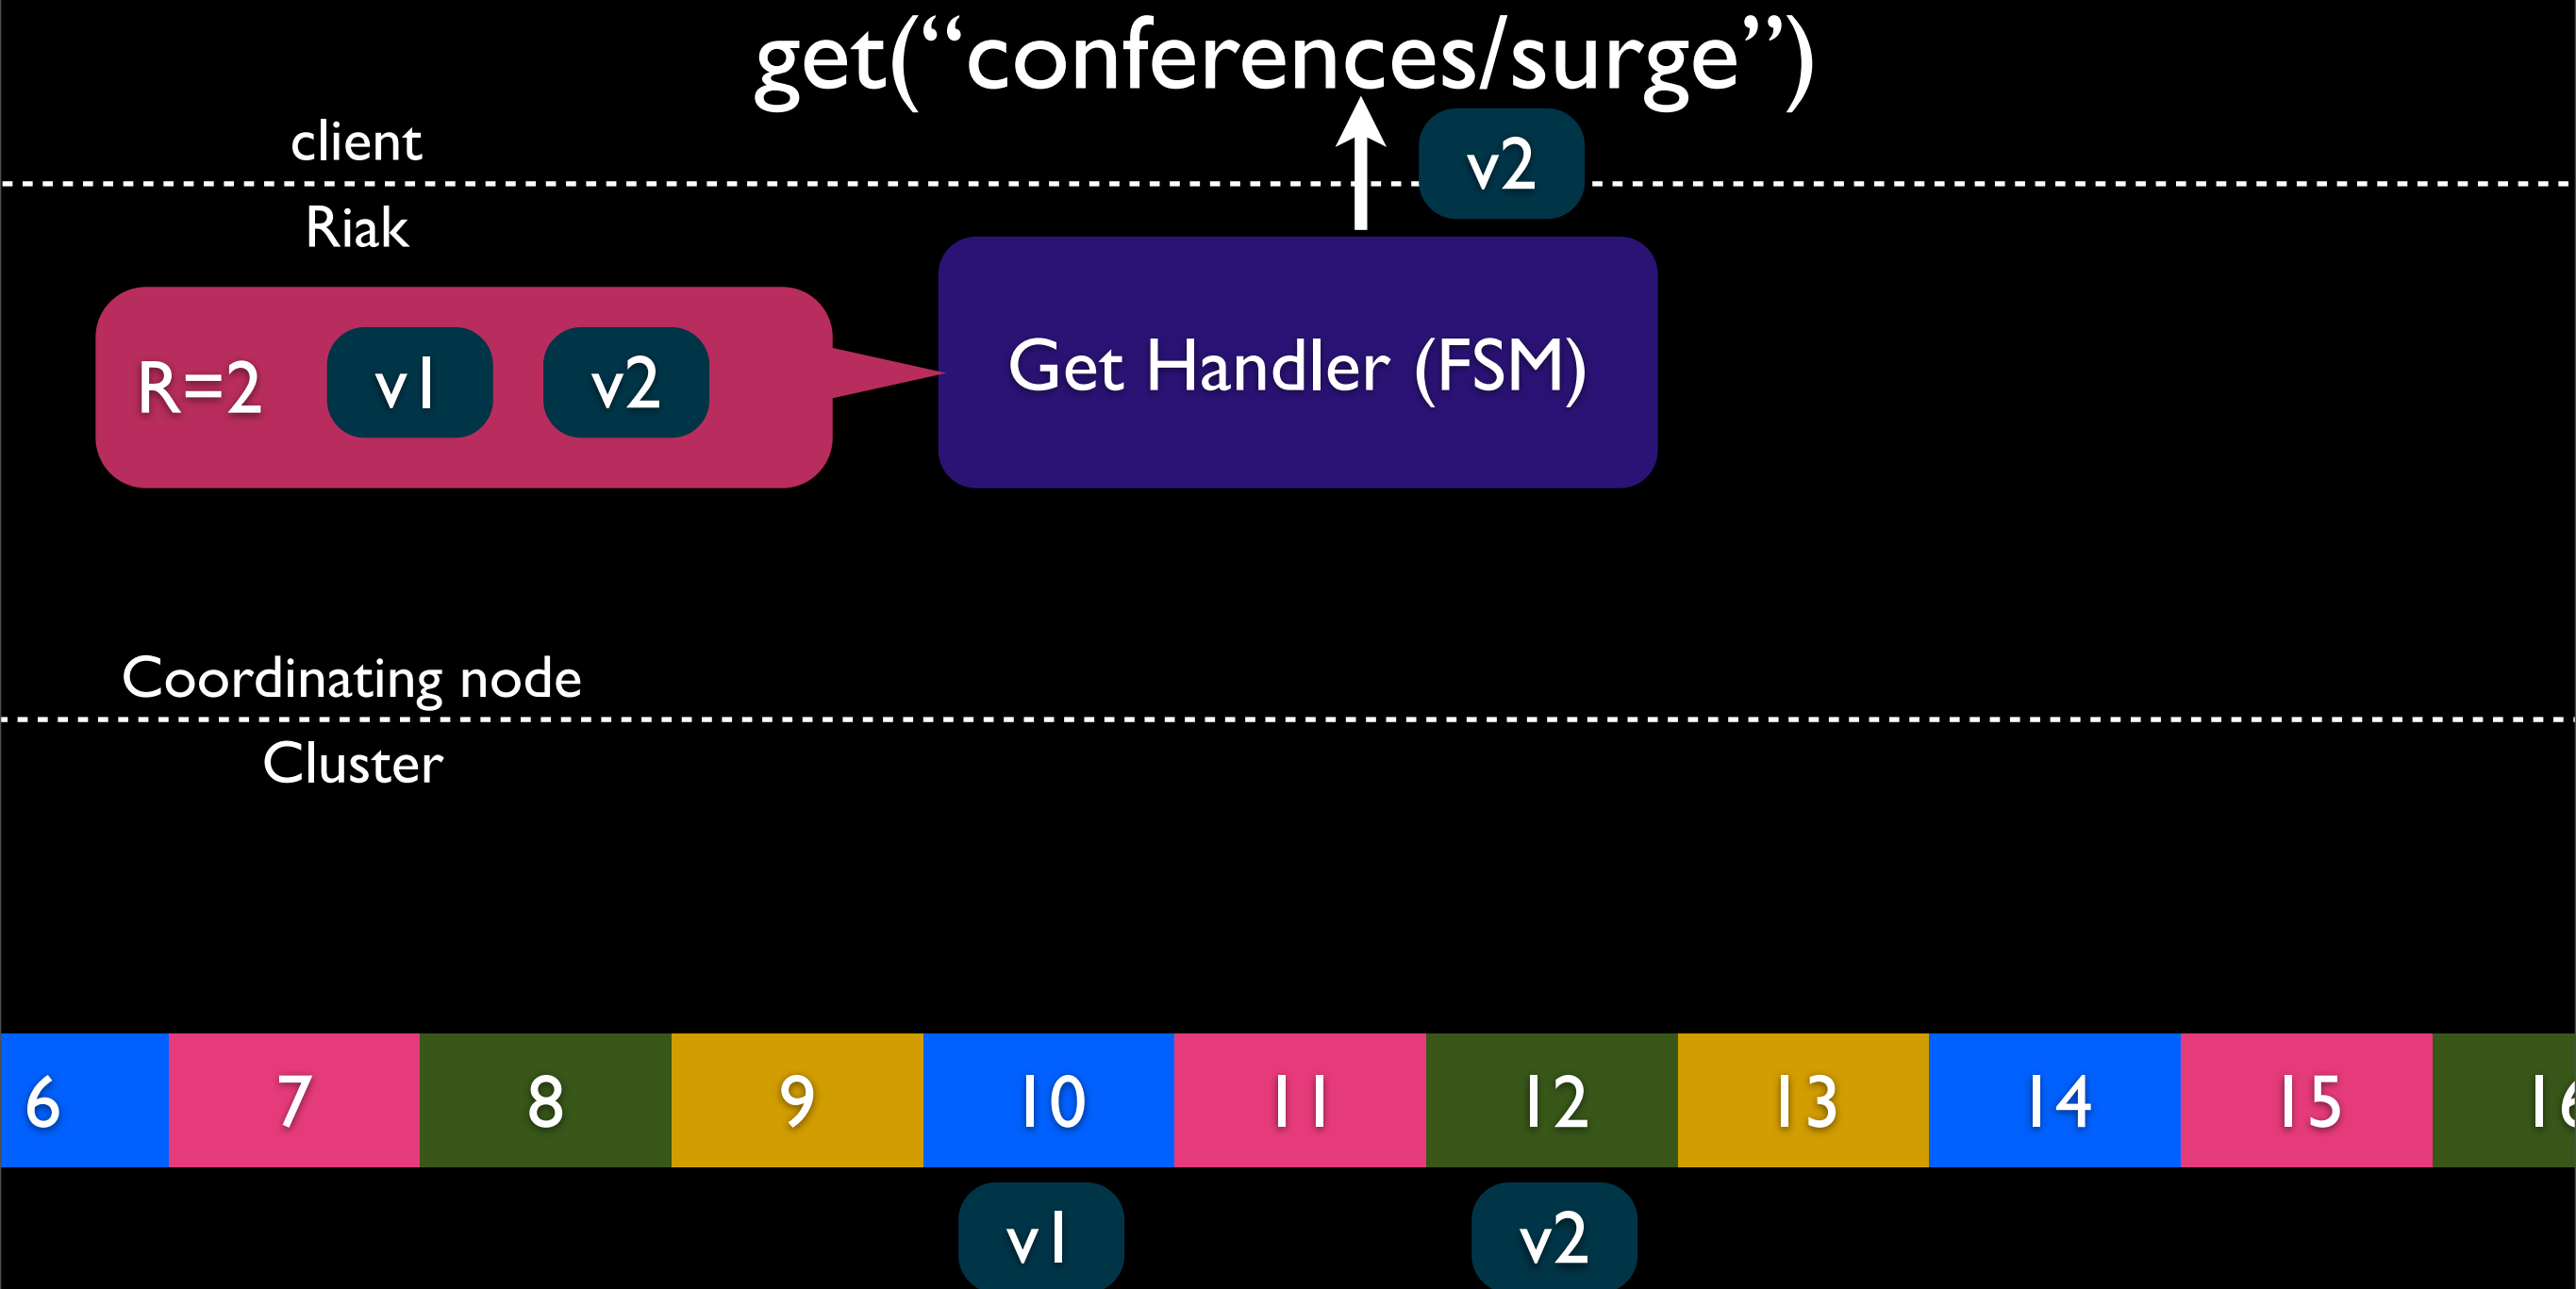


# Anatomy of a Request

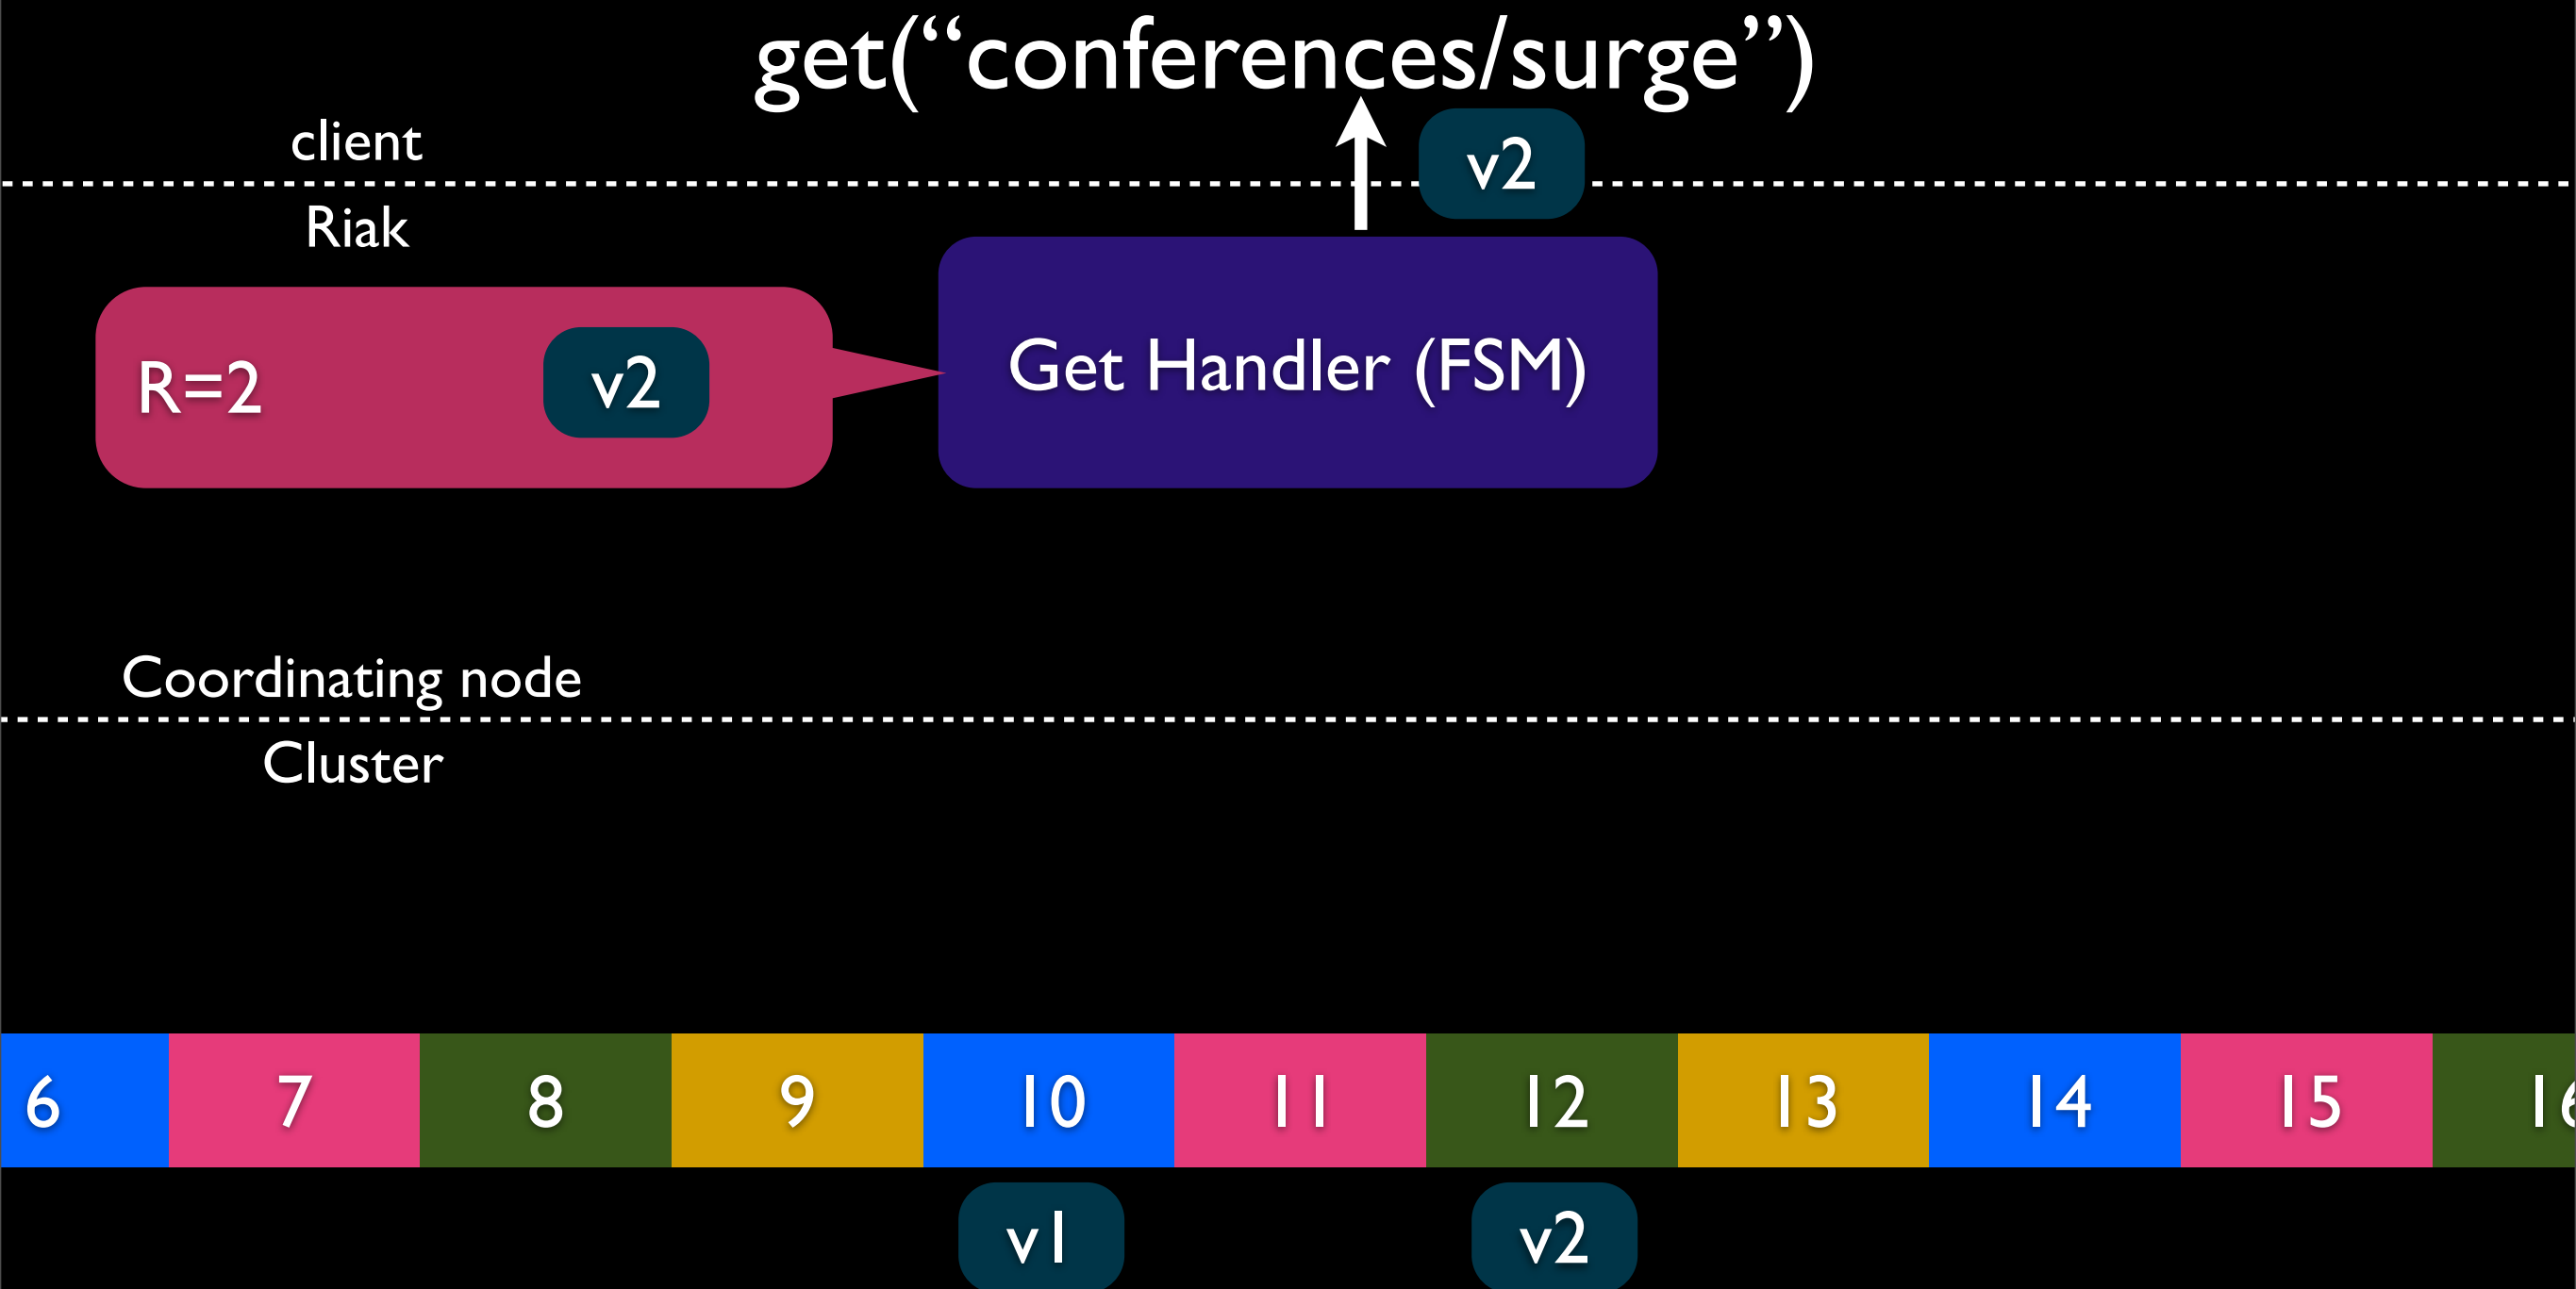
get("conferences/surge")

v2

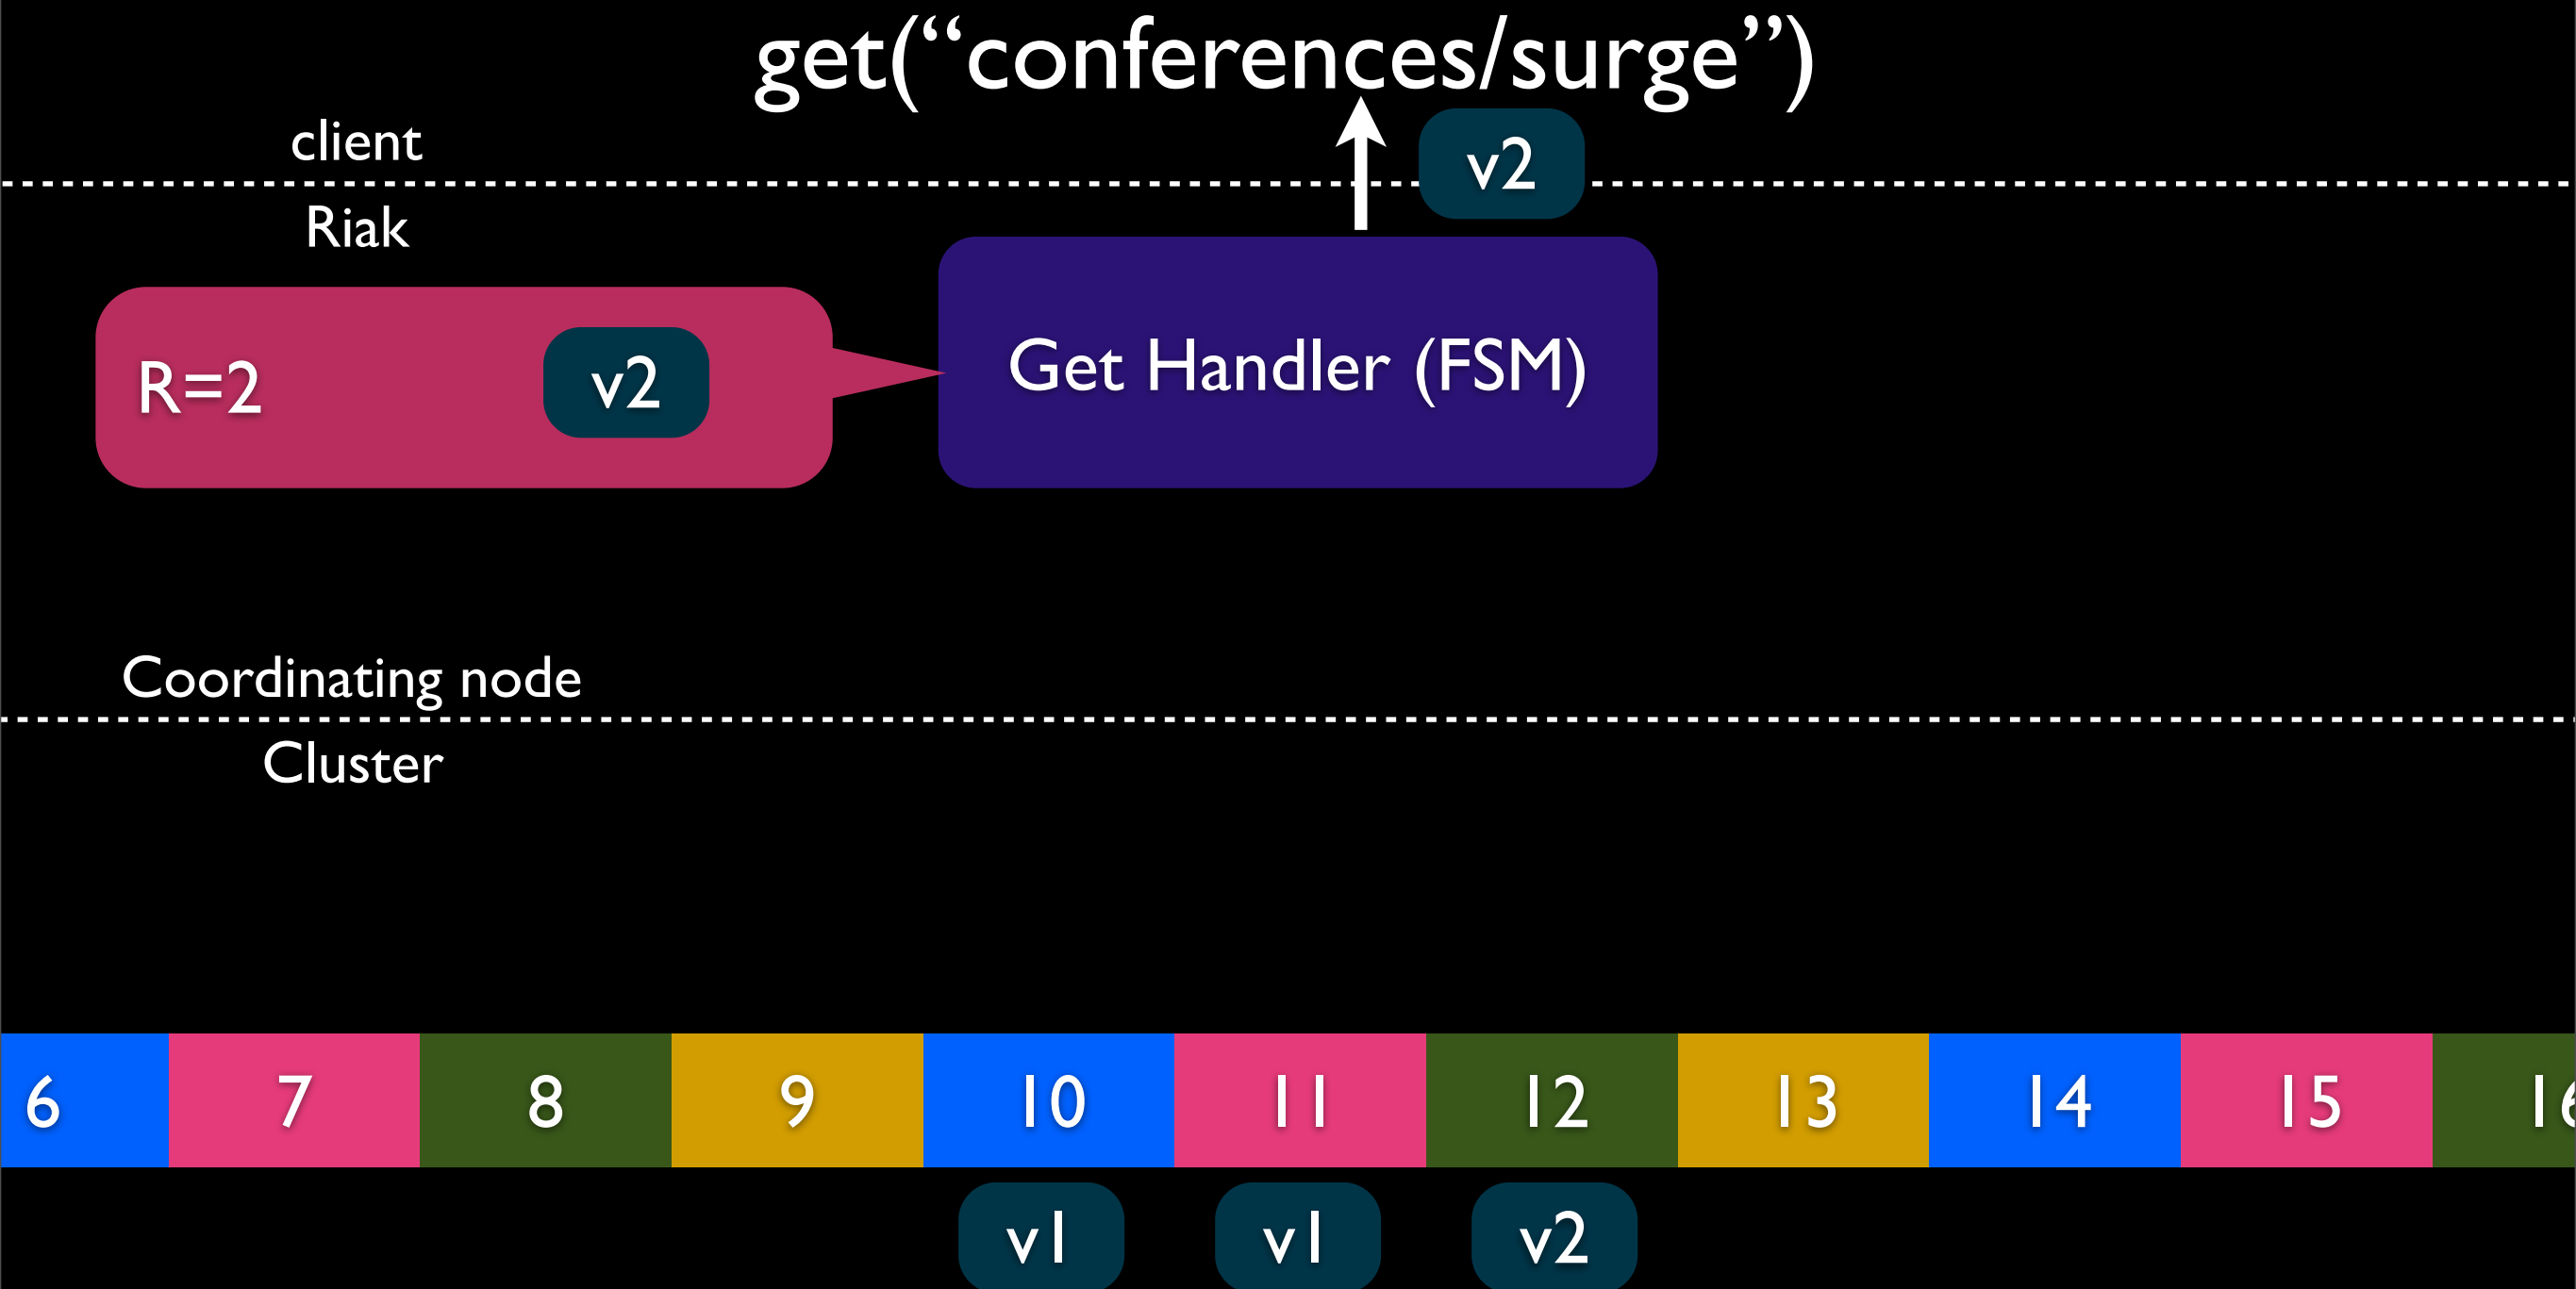
# Read Repair



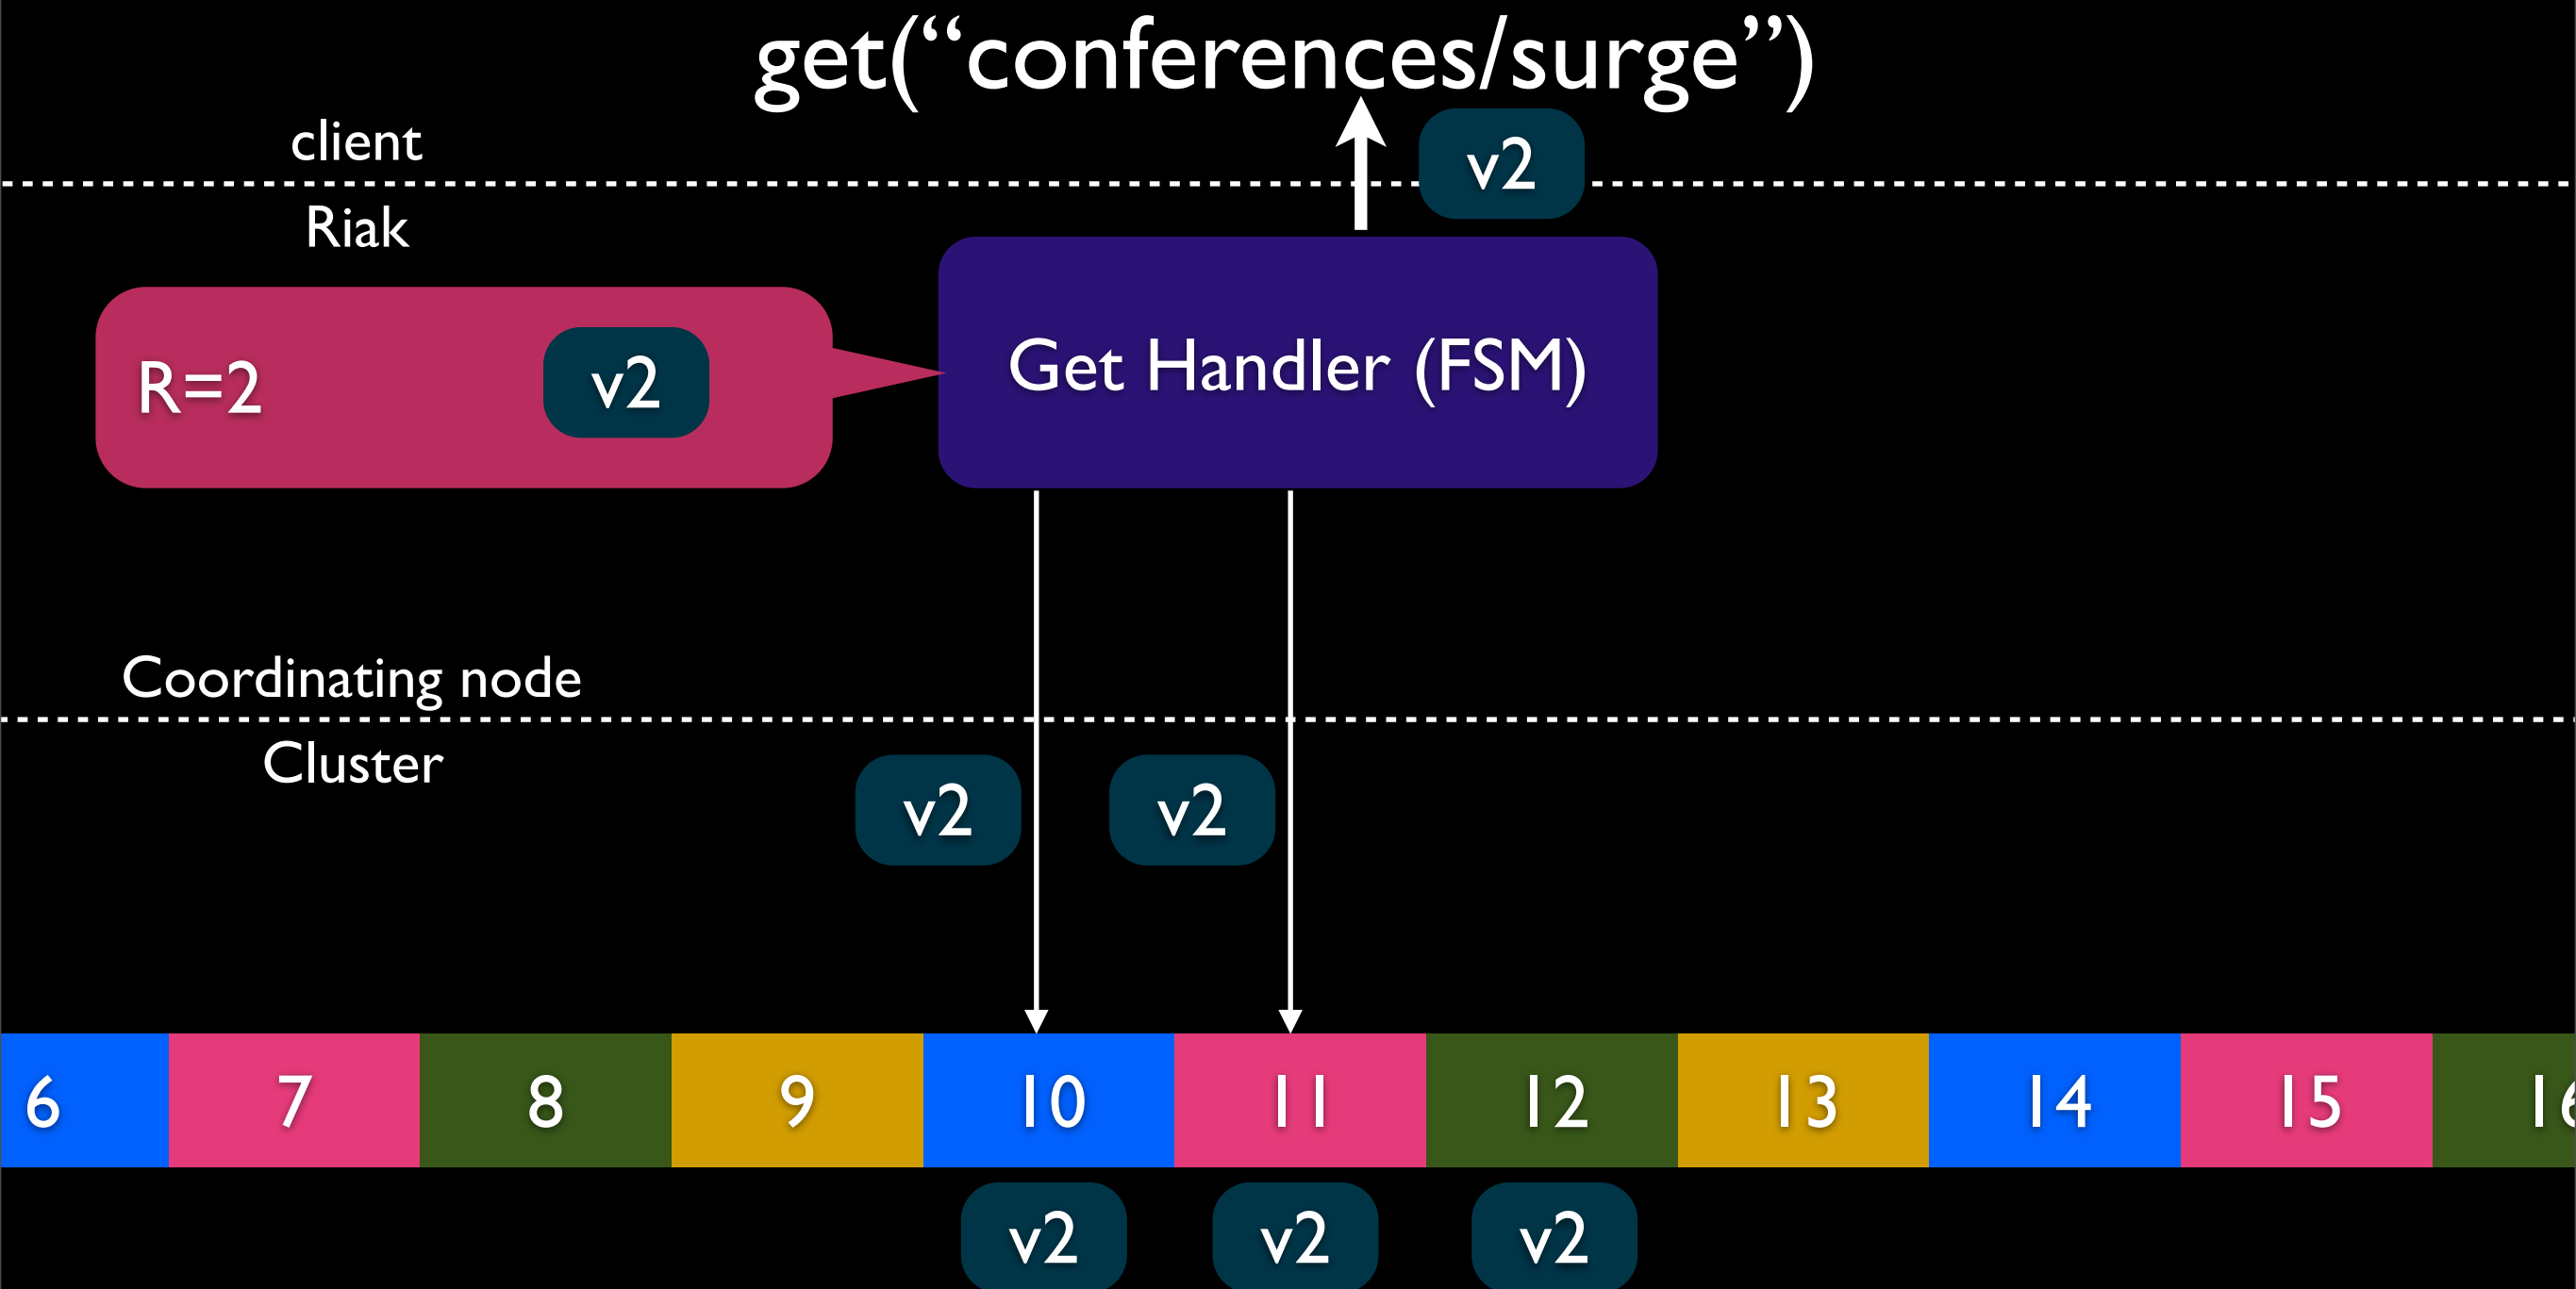
# Read Repair



# Read Repair



# Read Repair




# Riak Architecture

Erlang/OTP Runtime

# Riak Architecture

Erlang/OTP Runtime



Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

Riak KV



# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

Riak Core

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing

Riak Core

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing  
membership

Riak Core

Riak KV



# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing  
membership

handoff

Riak Core

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff  
membership   node-liveness

Riak Core

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff   gossip  
membership   node-liveness

Riak Core

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff   gossip  
membership   node-liveness   buckets

Riak Core

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff   gossip  
membership   node-liveness   buckets

Riak Core

vnode master

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff   gossip  
membership   node-liveness   buckets

Riak Core

vnode master

vnodes

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff   gossip  
membership   node-liveness   buckets

Riak Core

vnode master

vnodes

storage backend

Riak KV

# Riak Architecture

Erlang/OTP Runtime

Client APIs

HTTP

Protocol Buffers

Erlang local client

Request Coordination

get

put

delete

map-reduce

consistent hashing   handoff   gossip  
membership   node-liveness   buckets

Riak Core

vnode master

vnodes

storage backend

Workers

Riak KV



# Modeling & Querying

# Application Design

# Application Design

- No intrinsic schema

# Application Design

- No intrinsic schema
- Your application defines the structure and semantics

# Application Design

- No intrinsic schema
- Your application defines the structure and semantics
- Your application resolves conflicts (if you care)

# Modeling Tools

# Modeling Tools

- Key-Value

# Modeling Tools

- Key-Value
- Links



# Modeling Tools

- Key-Value
- Links
- Full-text Search

# Modeling Tools

- Key-Value
- Links
- Full-text Search
- Secondary Indexes (2I)

# Modeling Tools

- Key-Value
- Links
- Full-text Search
- Secondary Indexes (2I)
- MapReduce

# Key-Value

# Key-Value

- Content-Types

# Key-Value

- Content-Types
- Denormalize

# Key-Value

- Content-Types
- Denormalize
- Meaningful or “guessable” keys

# Key-Value

- Content-Types
- Denormalize
- Meaningful or “guessable” keys
  - Composites



# Key-Value

- Content-Types
- Denormalize
- Meaningful or “guessable” keys
  - Composites
  - Time-boxing

# Key-Value

- Content-Types
- Denormalize
- Meaningful or “guessable” keys
  - Composites
  - Time-boxing
- References (value is a key or list)

# Links

# Links

- Lightweight relationships, like <a>

# Links

- Lightweight relationships, like <a>
- Includes a “tag”

# Links

- Lightweight relationships, like <a>
- Includes a “tag”
- Built-in traversal op (“walking”)  
GET /riak/b/k/[bucket],[tag],[keep]

# Links

- Lightweight relationships, like <a>
- Includes a “tag”
- Built-in traversal op (“walking”)  
GET /riak/b/k/[bucket],[tag],[keep]
- Limited in number (part of meta)

# Full-text Search



# Full-text Search

- Designed for searching prose

# Full-text Search

- Designed for searching prose
- Lucene/Solr-like query interface

# Full-text Search

- Designed for searching prose
- Lucene/Solr-like query interface
- Automatically index K/V values

# Full-text Search

- Designed for searching prose
- Lucene/Solr-like query interface
- Automatically index K/V values
- Input to MapReduce

# Full-text Search

- Designed for searching prose
- Lucene/Solr-like query interface
- Automatically index K/V values
- Input to MapReduce
- Customizable index schemas

# Secondary Indexes

# Secondary Indexes

- Defined as metadata

# Secondary Indexes

- Defined as metadata
- Two index types: `_int` and `_bin`



# Secondary Indexes

- Defined as metadata
- Two index types: `_int` and `_bin`
- Two query types: `equal` and `range`

# Secondary Indexes

- Defined as metadata
- Two index types: `_int` and `_bin`
- Two query types: `equal` and `range`
- Input to MapReduce

# MapReduce

# MapReduce

- For more involved queries

# MapReduce

- For more involved queries
  - Specify input keys

# MapReduce

- For more involved queries
  - Specify input keys
  - Process data in “map” and “reduce” functions

# MapReduce

- For more involved queries
  - Specify input keys
  - Process data in “map” and “reduce” functions
    - JavaScript or Erlang

# MapReduce

- For more involved queries
  - Specify input keys
  - Process data in “map” and “reduce” functions
    - JavaScript or Erlang
- Not tuned for batch processing



# Lab: Querying

# Lab: Walk Links

# Lab: Walk Links

- Store an object with a Link

# Lab: Walk Links

- Store an object with a Link
- Store the target of the Link

# Lab: Walk Links

- Store an object with a Link
- Store the target of the Link
- Walk from one to the other

# Lab: Store Links

# Lab: Store Links

```
$ ./store-linked.sh
```

# Lab: Store Links

```
$ ./store-linked.sh  
Enter the origin key: sean  
Enter the target key: ian  
Enter the link's tag: coworker
```



# Lab: Store Links


```
$ ./store-linked.sh
Enter the origin key: sean
Enter the target key: ian
Enter the link's tag: coworker

*** Storing the origin ***
* About to connect() to localhost port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/sean HTTP/1.1
> Link: </riak/surge/ian>;riaktag="coworker"
. . .
< HTTP/1.1 204 No Content
. . .
```

# Lab: Store Links

```
$ ./store-linked.sh
Enter the origin key: sean
Enter the target key: ian
Enter the link's tag: coworker


*** Storing the origin ***
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/sean HTTP/1.1
> Link: </riak/surge/ian>;riaktag="coworker"
...
< HTTP/1.1 204 No Content
...
```



# Lab: Store Links

```
$ ./store-linked.sh
Enter the origin key: sean
Enter the target key: ian
Enter the link's tag: coworker


*** Storing the origin ***
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/sean HTTP/1.1
> Link: </riak/surge/ian>;riaktag="coworker"
...
< HTTP/1.1 204 No Content
...
```



# Lab: Store Links

```
$ ./store-linked.sh
Enter the origin key: sean
Enter the target key: ian
Enter the link's tag: coworker

*** Storing the origin ***
* About to connect() to localhost port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/sean HTTP/1.1
> Link: </riak/surge/ian>;riaktag="coworker"
...
< HTTP/1.1 204 No Content
...
```



# Lab: Store Links

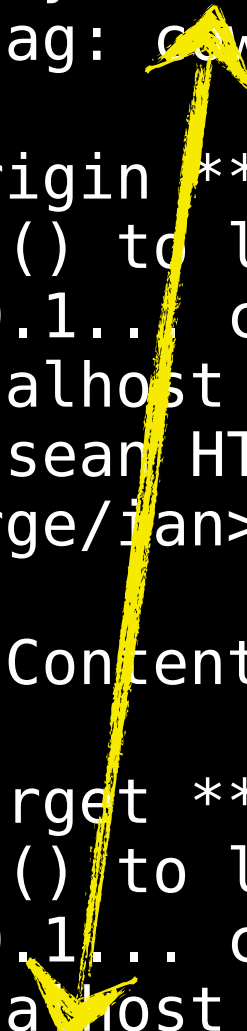
```
$ ./store-linked.sh
Enter the origin key: sean
Enter the target key: ian
Enter the link's tag: coworker

*** Storing the origin ***
* About to connect() to localhost port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/sean HTTP/1.1
> Link: </riak/surge/ian>;riaktag="coworker"
. . .
< HTTP/1.1 204 No Content
. . .
*** Storing the target ***
* About to connect() to localhost port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/ian HTTP/1.1
. . .
< HTTP/1.1 204 No Content
```

# Lab: Store Links

```
$ ./store-linked.sh
Enter the origin key: sean
Enter the target key: ian
Enter the link's tag: coworker

*** Storing the origin ***
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/sean HTTP/1.1
> Link: </riak/surge/ian>;riaktag="coworker"
...
< HTTP/1.1 204 No Content
...
*** Storing the target ***
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/ian HTTP/1.1
...
< HTTP/1.1 204 No Content
```



# Lab: Walk Links (1/3)

# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
```



# Lab: Walk Links (1/3)


```
$ ./walk-linked.sh  
Enter the origin key: sean  
Enter the link spec (default: _,_,_):
```

# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
Enter the origin key: sean
Enter the link spec (default: _,_,_):
...
> GET /riak/surge/sean/_,_,_ HTTP/1.1
...
```


# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh  
Enter the origin key: sean  
Enter the link spec (default: _,_,_):  
...  
> GET /riak/surge/sean/_,_,_ HTTP/1.1  
...
```



# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
Enter the origin key: sean
Enter the link spec (default: _,_,_):
...
> GET /riak/surge/sean/_,_,_ HTTP/1.1
...
```



# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
Enter the origin key: sean
Enter the link spec (default: _,_,_):
...
> GET /riak/surge/sean/_,_,_ HTTP/1.1
...
< HTTP/1.1 200 OK
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Expires: Tue, 27 Sep 2011 20:51:37 GMT
< Date: Tue, 27 Sep 2011 20:41:37 GMT
< Content-Type: multipart/mixed; boundary=J602SZfLfSEdAv5dv3mttR7AV0F
< Content-Length: 438
<
```

# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
Enter the origin key: sean
Enter the link spec (default: _,_,_):
...
> GET /riak/surge/sean/_,_,_ HTTP/1.1
...
< HTTP/1.1 200 OK
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Expires: Tue, 27 Sep 2011 20:51:37 GMT
< Date: Tue, 27 Sep 2011 20:41:37 GMT
< Content-Type: multipart/mixed; boundary=J602SZfLfSEdAv5dv3mttR7AV0F
< Content-Length: 438
<
```

# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
Enter the origin key: sean
Enter the link spec (default: _,_,_):
...
> GET /riak/surge/sean/_,_,_ HTTP/1.1
...
< HTTP/1.1 200 OK
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Expires: Tue, 27 Sep 2011 20:51:37 GMT
< Date: Tue, 27 Sep 2011 20:41:37 GMT
< Content-Type: multipart/mixed; boundary=J602SZfLfSEdAv5dv3mttR7AV0F
< Content-Length: 438
<

--J602SZfLfSEdAv5dv3mttR7AV0F
Content-Type: multipart/mixed; boundary=C1NUMpwbmSvb7CbqEAY1KdYRiUH

--C1NUMpwbmSvb7CbqEAY1KdYRiUH
X-Riak-Vclock: a85hYGBgzGDKBVIcMRuuc/k1GU/KYEpkymNl0Nzw7ThfFgA=
Location: /riak/surge/ian
Content-Type: text/plain
Link: </riak/surge>; rel="up"
Etag: 51h3q7RjTNaHWYp04P0MJj
Last-Modified: Tue, 27 Sep 2011 20:38:01 GMT

target
--C1NUMpwbmSvb7CbqEAY1KdYRiUH--

--J602SZfLfSEdAv5dv3mttR7AV0F--
```

# Lab: Walk Links (1/3)

```
$ ./walk-linked.sh
Enter the origin key: sean
Enter the link spec (default: _,_,_):
...
> GET /riak/surge/sean/_,_,_ HTTP/1.1
...
< HTTP/1.1 200 OK
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Expires: Tue, 27 Sep 2011 20:51:37 GMT
< Date: Tue, 27 Sep 2011 20:41:37 GMT
< Content-Type: multipart/mixed; boundary=J602SZfLfSEdAv5dv3mttR7AV0F
< Content-Length: 438
<

--J602SZfLfSEdAv5dv3mttR7AV0F
Content-Type: multipart/mixed; boundary=C1NUMpwbmSvb7CbqEAY1KdYRiUH

--C1NUMpwbmSvb7CbqEAY1KdYRiUH
X-Riak-Verlock: a85hYGBgzGDKBVIcMRuuc/k1GU/KYEpkymNl0Nzw7ThfFgA=
Location: /riak/surge/ian
Content-Type: text/plain
Link: </riak/surge>; rel="up"
Etag: 51h3q7RjTNaHWYp04P0MJj
Last-Modified: Tue, 27 Sep 2011 20:38:01 GMT

target
--C1NUMpwbmSvb7CbqEAY1KdYRiUH--

--J602SZfLfSEdAv5dv3mttR7AV0F--
```



# Lab: Walk Links (2/3)

# Lab: Walk Links (2/3)

Enter the origin key: sean

Enter the link spec (default: \_,\_,\_): surge,\_,\_

# Lab: Walk Links (2/3)

Enter the origin key: sean

Enter the link spec (default: \_,\_,\_): surge,\_,\_

> GET /riak/surge/sean/surge,\_,\_ HTTP/1.1

...

# Lab: Walk Links (2/3)

Enter the origin key: sean

Enter the link spec (default: \_,\_,\_): surge,\_,\_

```
> GET /riak/surge/sean/surge,_,_ HTTP/1.1
```

```
...
```

```
>
```

```
< HTTP/1.1 200 OK
```

```
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
```

```
< Expires: Tue, 27 Sep 2011 20:52:53 GMT
```

```
< Date: Tue, 27 Sep 2011 20:42:53 GMT
```

```
< Content-Type: multipart/mixed; boundary=CsxMIs09tfadrRJ7EQ3XL2ivQ4f
```

```
< Content-Length: 438
```

```
<
```

```
--CsxMIs09tfadrRJ7EQ3XL2ivQ4f
```

```
Content-Type: multipart/mixed; boundary=I3U33LkzkqJi6HtbwsU0qRd4k4y
```

```
--I3U33LkzkqJi6HtbwsU0qRd4k4y
```

```
X-Riak-Vclock: a85hYGBgzGDKBVicMRuuc/k1GU/KYEpkymNl0Nzw7ThfFgA=
```

```
Location: /riak/surge/sean
```

```
Content-Type: text/plain
```

```
Link: </riak/surge>; rel="up"
```

```
Etag: 51h3q7RjTNaHWYp04P0MJj
```

```
Last-Modified: Tue, 27 Sep 2011 20:38:01 GMT
```

```
target
```

```
--I3U33LkzkqJi6HtbwsU0qRd4k4y--
```

```
--CsxMIs09tfadrRJ7EQ3XL2ivQ4f--
```

# Lab: Walk Links (3/3)

# Lab: Walk Links (3/3)

```
Enter the origin key: sean  
Enter the link spec (default: _,_,_): foo,_,_
```

# Lab: Walk Links (3/3)

Enter the origin key: sean

Enter the link spec (default: \_\_,\_\_,\_\_): foo,\_\_,\_\_

...

```
> GET /riak/surge/sean/foo,__,__ HTTP/1.1
```

```
< HTTP/1.1 200 OK
```

```
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
```

```
< Expires: Tue, 27 Sep 2011 20:53:08 GMT
```

```
< Date: Tue, 27 Sep 2011 20:43:08 GMT
```

```
< Content-Type: multipart/mixed; boundary=3AsxaHhVMlDEQCakLSmNqQUTS4Y
```

```
< Content-Length: 172
```

```
<
```

```
--3AsxaHhVMlDEQCakLSmNqQUTS4Y
```

```
Content-Type: multipart/mixed; boundary=Ivtq9RgHpydECNEZ0nHpiqHcFYl
```

```
--Ivtq9RgHpydECNEZ0nHpiqHcFYl--
```

```
--3AsxaHhVMlDEQCakLSmNqQUTS4Y--
```

# Lab: Secondary Indexes



# Lab: Secondary Indexes

- Create some objects with indexes

# Lab: Secondary Indexes

- Create some objects with indexes
- Query the index for their keys

# Lab: Secondary Indexes

- Create some objects with indexes
- Query the index for their keys
- Query input to MapReduce

# Lab: Create 21 Objects

# Lab: Create 21 Objects

```
$ ./store-indexed.sh
```

# Lab: Create 21 Objects

```
$ ./store-indexed.sh  
Enter the number of objects to create: 40
```

# Lab: Create 21 Objects

```
$ ./store-indexed.sh
Enter the number of objects to create: 40
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/surgeobj1?returnbody=true HTTP/1.1
> User-Agent: curl/7.19.7 (universal-apple-darwin10.0) libcurl/7.19.7 OpenSSL/0.9.8r
zlib/1.2.3
> Host: localhost:8091
> Accept: */*
> Content-Type: application/json
> X-Riak-Index-surgeobj_int: 1
> Content-Length: 14
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1uZVkMCUy5rEyHDn57ThfFgA=
< x-riak-index-surgeobj_int: 1
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Date: Tue, 27 Sep 2011 22:27:16 GMT
< Content-Type: application/json
< Content-Length: 14
<
* Connection #0 to host localhost left intact
* Closing connection #0
{"surgeobj":1} ...
```

# Lab: Create 21 Objects

```
$ ./store-indexed.sh
Enter the number of objects to create: 40
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/surgeobj1?returnbody=true HTTP/1.1
> User-Agent: curl/7.19.7 (universal-apple-darwin10.0) libcurl/7.19.7 OpenSSL/0.9.8r
zlib/1.2.3
> Host: localhost:8091
> Accept: */*
> Content-Type: application/json
> X-Riak-Index-surgeobj_int: 1
> Content-Length: 14
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVIcR4M2cvs1uZVkMCUy5rEyHDn57ThfFgA=
< x-riak-index-surgeobj_int: 1
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Date: Tue, 27 Sep 2011 22:27:16 GMT
< Content-Type: application/json
< Content-Length: 14
<
* Connection #0 to host localhost left intact
* Closing connection #0
{"surgeobj":1} ...
```


key



# Lab: Create 21 Objects

```
$ ./store-indexed.sh
Enter the number of objects to create: 40
* About to connect() to localhost port 8091 (#0)
* Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> PUT /riak/surge/surgeobj1?returnbody=true HTTP/1.1
> User-Agent: curl/7.19.7 (universal-apple-darwin10.0) libcurl/7.19.7 OpenSSL/0.9.8r
zlib/1.2.3
> Host: localhost:8091
> Accept: */*
> Content-Type: application/json
> X-Riak-Index-surgeobj_int: 1
> Content-Length: 14
>
< HTTP/1.1 200 OK
< X-Riak-Vclock: a85hYGBgzGDKBVTcPm2cvs1uZVkMCUy5rEyHDn57ThfFgA=
< x-riak-index-surgeobj_int: 1
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Link: </riak/surge>; rel="up"
< Date: Tue, 27 Sep 2011 22:27:16 GMT
< Content-Type: application/json
< Content-Length: 14
<
* Connection #0 to host localhost left intact
* Closing connection #0
{"surgeobj":1} ...
```

**index header**



# Lab: Query 2I (1/2)

# Lab: Query 2I (1/2)

```
$ ./query-indexed.sh
```

# Lab: Query 2I (1/2)

```
$ ./query-indexed.sh  
Query on range or equality? (r/e) e  
Index value: 10
```

# Lab: Query 2I (1/2)

```
$ ./query-indexed.sh
Query on range or equality? (r/e) e
Index value: 10

* About to connect() to localhost port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> GET /buckets/surge/index/surgeobj_int/10 HTTP/1.1
> User-Agent: curl/7.19.7 (universal-apple-darwin10.0) libcurl/7.19.7
  OpenSSL/0.9.8r zlib/1.2.3
> Host: localhost:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Date: Tue, 27 Sep 2011 22:27:27 GMT
< Content-Type: application/json
< Content-Length: 23
<
* Connection #0 to host localhost left intact
* Closing connection #0
{"keys":["surgeobj10"]}
```

# Lab: Query 2I (2/2)

# Lab: Query 2I (2/2)

Query on range or equality? (r/e) r

Range start: 5

Range end: 15

# Lab: Query 2I (2/2)

Query on range or equality? (r/e) r

Range start: 5

Range end: 15

```
* About to connect() to localhost port 8091 (#0)
*   Trying 127.0.0.1... connected
* Connected to localhost (127.0.0.1) port 8091 (#0)
> GET /buckets/surge/index/surgeobj_int/5/15 HTTP/1.1
> User-Agent: curl/7.19.7 (universal-apple-darwin10.0) libcurl/7.19.7
  OpenSSL/0.9.8r zlib/1.2.3
> Host: localhost:8091
> Accept: */*
>
< HTTP/1.1 200 OK
< Vary: Accept-Encoding
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Date: Tue, 27 Sep 2011 22:27:36 GMT
< Content-Type: application/json
< Content-Length: 148
<
* Connection #0 to host localhost left intact
* Closing connection #0
{"keys":
["surgeobj12","surgeobj15","surgeobj6","surgeobj5","surgeobj13","surgeobj8",
"surgeobj11","surgeobj14","surgeobj9","surgeobj10","surgeobj7"]}
```



# Lab: 21 MapReduce

# Lab: 21 MapReduce

```
$ ./mapred-indexed.sh
```

# Lab: 21 MapReduce

```
$ ./mapred-indexed.sh  
Query on range or equality? (r/e) r  
Range start: 35  
Range end: 40
```

# Lab: 21 MapReduce

```
$ ./mapred-indexed.sh  
Query on range or equality? (r/e) r  
Range start: 35  
Range end: 40
```

```
Executing MapReduce:  
{ "inputs":{ "bucket":"surge", "index":"surgeobj_int" , "start":35, "end":40},  
  "query":[ {"map":  
{"name":"Riak.mapValuesJson", "language":"javascript", "keep":true}} ]}
```

# Lab: 21 MapReduce

```
$ ./mapred-indexed.sh
Query on range or equality? (r/e) r
Range start: 35
Range end: 40
```

Executing MapReduce:

```
{ "inputs":{ "bucket":"surge", "index":"surgeobj_int" , "start":35, "end":40},
  "query":[ { "map":
{"name":"Riak.mapValuesJson", "language":"javascript", "keep":true}} ]}
```

```
> POST /mapred HTTP/1.1
> User-Agent: curl/7.19.7 (universal-apple-darwin10.0) libcurl/7.19.7 OpenSSL/
0.9.8r zlib/1.2.3
> Host: localhost:8091
> Accept: */*
> Content-Type: application/json
> Content-Length: 164
>
< HTTP/1.1 200 OK
< Server: MochiWeb/1.1 WebMachine/1.9.0 (participate in the frantic)
< Date: Wed, 28 Sep 2011 01:41:53 GMT
< Content-Type: application/json
< Content-Length: 97
<
* Connection #0 to host localhost left intact
* Closing connection #0
[{"surgeobj":35}, {"surgeobj":40}, {"surgeobj":36}, {"surgeobj":37}, {"surgeobj":
39}, {"surgeobj":38}]
```

# Riak Operations

# Configuration

# File Locations



# File Locations

- Configuration

# File Locations

- Configuration
  - /etc/riak

# File Locations

- Configuration
  - /etc/riak
- Binaries

# File Locations

- Configuration
  - `/etc/riak`
- Binaries
  - `/usr/sbin/riak`

# File Locations

- Configuration
  - /etc/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin

# File Locations

- Configuration
  - `/etc/riak`
- Binaries
  - `/usr/sbin/riak`
  - `/usr/sbin/riak-admin`
  - `/usr/[lib,lib64]/riak`

# File Locations

- Configuration
  - /etc/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin
  - /usr/[lib,lib64]/riak
- Logs

# File Locations

- Configuration
  - /etc/riak
- Logs
  - /var/log/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin
  - /usr/[lib,lib64]/riak



# File Locations

- Configuration
  - /etc/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin
  - /usr/[lib,lib64]/riak
- Logs
  - /var/log/riak
- Data

# File Locations

- Configuration
  - /etc/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin
  - /usr/[lib,lib64]/riak
- Logs
  - /var/log/riak
- Data
  - /var/lib/riak

# File Locations

- Configuration
  - /etc/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin
  - /usr/[lib,lib64]/riak
- Logs
  - /var/log/riak
- Data
  - /var/lib/riak
- Handles / Temp Files

# File Locations

- Configuration
  - /etc/riak
- Binaries
  - /usr/sbin/riak
  - /usr/sbin/riak-admin
  - /usr/[lib,lib64]/riak
- Logs
  - /var/log/riak
- Data
  - /var/lib/riak
- Handles / Temp Files
  - /tmp/riak

**/etc/riak/vm.args**

# /etc/riak/vm.args

- Overview

# /etc/riak/vm.args

- Overview
  - Erlang VM configuration settings.

# /etc/riak/vm.args

- Overview
  - Erlang VM configuration settings.
- Important Settings



# /etc/riak/vm.args

- Overview
  - Erlang VM configuration settings.
- Important Settings
  - Node name.

# /etc/riak/vm.args

- Overview
  - Erlang VM configuration settings.
- Important Settings
  - Node name.
  - Security cookie.

# /etc/riak/vm.args

- Overview
  - Erlang VM configuration settings.
- Important Settings
  - Node name.
  - Security cookie.
- Documentation

# /etc/riak/vm.args

- Overview
  - Erlang VM configuration settings.
- Important Settings
  - Node name.
  - Security cookie.
- Documentation
  - <http://www.erlang.org/doc/man/erl.html>

**/etc/riak/vm.args**

# /etc/riak/vm.args

- Node Name & Security Cookie  
## Name of the riak node

# /etc/riak/vm.args

- Node Name & Security Cookie

## Name of the riak node

-name riak@127.0.0.1

# /etc/riak/vm.args

- Node Name & Security Cookie

## Name of the riak node

-name riak@127.0.0.1

## Cookie for distributed erlang



# /etc/riak/vm.args

- Node Name & Security Cookie

```
## Name of the riak node
```

```
-name riak@127.0.0.1
```

```
## Cookie for distributed erlang
```

```
-setcookie riak
```

**/etc/riak/app.config**

# /etc/riak/app.config

- Overview

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports
  - Directories

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports
  - Directories
  - Number of Partitions



# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports
  - Directories
  - Number of Partitions
  - Storage Backend

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports
  - Directories
  - Number of Partitions
  - Storage Backend
- Documentation

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports
  - Directories
  - Number of Partitions
  - Storage Backend
- Documentation
  - <http://wiki.basho.com/Configuration-Files.html>

# /etc/riak/app.config

- Overview
  - Riak (and dependency) configuration settings.
- Important Settings
  - Ports
  - Directories
  - Number of Partitions
  - Storage Backend
- Documentation
  - <http://wiki.basho.com/Configuration-Files.html>
  - Inline comments

**/etc/riak/app.config**

# /etc/riak/app.config

- Storage Engine

%% Storage\_backend specifies the Erlang module defining the storage

# /etc/riak/app.config

- Storage Engine

%% Storage\_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.

# /etc/riak/app.config

- Storage Engine

```
%% Storage_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.  
{storage_backend, riak_kv_bitcask_backend},
```



# /etc/riak/app.config

- Storage Engine

%% Storage\_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.  
{storage\_backend, riak\_kv\_bitcask\_backend},

- About Storage Engines

# /etc/riak/app.config

- Storage Engine

%% Storage\_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.  
{storage\_backend, riak\_kv\_bitcask\_backend},

- About Storage Engines

- Riak has pluggable storage engines. (Bitcask, Innostore, LevelDB)

# /etc/riak/app.config

- Storage Engine

```
%% Storage_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.  
{storage_backend, riak_kv_bitcask_backend},
```

- About Storage Engines

- Riak has pluggable storage engines. (Bitcask, Innostore, LevelDB)
- Different engines have different tradeoffs.

# /etc/riak/app.config

- Storage Engine

```
%% Storage_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.  
{storage_backend, riak_kv_bitcask_backend},
```

- About Storage Engines

- Riak has pluggable storage engines. (Bitcask, Innostore, LevelDB)
- Different engines have different tradeoffs.
- Engine selection depends on shape of data.

# /etc/riak/app.config

- Storage Engine

```
%% Storage_backend specifies the Erlang module defining the storage  
%% mechanism that will be used on this node.  
{storage_backend, riak_kv_bitcask_backend},
```

- About Storage Engines

- Riak has pluggable storage engines. (Bitcask, Innostore, LevelDB)
- Different engines have different tradeoffs.
- Engine selection depends on shape of data.
- More on this later.

# /etc/riak/app.config

```
%% Bitcask Config
{bitcask, [
    {data_root, "data/bitcask"}
]},

{leveldb, [
    {data_root, "data/leveldb"}
]},
```

# /etc/riak/app.config

- File Locations

```
%% Bitcask Config
```

```
{bitcask, [  
    {data_root, "data/bitcask"}  
]},
```

```
{leveldb, [  
    {data_root, "data/leveldb"}  
]},
```

# /etc/riak/app.config

- File Locations

```
%% Bitcask Config
```

```
{bitcask, [  
    {data_root, "data/bitcask"}  
]},
```

```
{leveldb, [  
    {data_root, "data/leveldb"}  
]},
```



# /etc/riak/app.config

- File Locations

```
%% Bitcask Config
```

```
{bitcask, [  
    {data_root, "data/bitcask"}  
]},
```

```
{leveldb, [  
    {data_root, "data/leveldb"}  
]},
```

# /etc/riak/app.config

- File Locations

```
%% Bitcask Config
```

```
{bitcask, [  
    {data_root, "data/bitcask"}  
]},
```

```
{leveldb, [  
    {data_root, "data/leveldb"}  
]},
```

# /etc/riak/app.config

- File Locations

```
%% Bitcask Config
```

```
{bitcask, [  
    {data_root, "data/bitcask"}  
]},
```

```
{leveldb, [  
    {data_root, "data/leveldb"}  
]},
```

# /etc/riak/app.config

```
%% http is a list of IP addresses and TCP ports that the  
Riak
```

```
%% HTTP interface will bind.
```

```
{http, [ {"127.0.0.1", 8091 } ]},
```

```
%% pb_ip is the IP address that the Riak Protocol Buffers  
interface
```

```
%% will bind to.  If this is undefined, the interface will  
not run.
```

```
{pb_ip,  "127.0.0.1" },
```

```
%% pb_port is the TCP port that the Riak Protocol Buffers  
interface
```

```
%% will bind to
```

```
{pb_port, 8081 },
```

# /etc/riak/app.config

- HTTP & Protocol Buffers

%% http is a list of IP addresses and TCP ports that the Riak

%% HTTP interface will bind.

```
{http, [ {"127.0.0.1", 8091 } ]},
```

%% pb\_ip is the IP address that the Riak Protocol Buffers interface

%% will bind to. If this is undefined, the interface will not run.

```
{pb_ip, "127.0.0.1" },
```

%% pb\_port is the TCP port that the Riak Protocol Buffers interface

%% will bind to

```
{pb_port, 8081 },
```

# Securing Riak

# Securing Riak

- Overview

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”



# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.
  - Riak assumes internal environment is trusted.

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.
  - Riak assumes internal environment is trusted.
- Securing the HTTP Interface

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.
  - Riak assumes internal environment is trusted.
- Securing the HTTP Interface
  - HTTP Auth via Proxy

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.
  - Riak assumes internal environment is trusted.
- Securing the HTTP Interface
  - HTTP Auth via Proxy
  - Does support SSL

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.
  - Riak assumes internal environment is trusted.
- Securing the HTTP Interface
  - HTTP Auth via Proxy
  - Does support SSL
- Securing Protocol Buffers

# Securing Riak

- Overview
  - Erlang has a “cookie” for “security.”
    - Do not rely on it. Plain text.
  - Riak assumes internal environment is trusted.
- Securing the HTTP Interface
  - HTTP Auth via Proxy
  - Does support SSL
- Securing Protocol Buffers
  - No security.

# Load Balancing



# Load Balancing

- Overview

# Load Balancing

- Overview
  - Any node can handle any request.

# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.

# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.
- Load Balancing HTTP

# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.
- Load Balancing HTTP
  - HAProxy, nginx, etc...

# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.
- Load Balancing HTTP
  - HAProxy, nginx, etc...
- Load Balancing Protocol Buffers

# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.
- Load Balancing HTTP
  - HAProxy, nginx, etc...
- Load Balancing Protocol Buffers
  - Any TCP Load Balancer

# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.
- Load Balancing HTTP
  - HAProxy, nginx, etc...
- Load Balancing Protocol Buffers
  - Any TCP Load Balancer
- In General



# Load Balancing

- Overview
  - Any node can handle any request.
  - Makes load balancing easy.
- Load Balancing HTTP
  - HAProxy, nginx, etc...
- Load Balancing Protocol Buffers
  - Any TCP Load Balancer
- In General
  - Use “least connected” strategy.

# Storage Backends

**There are 5 to choose  
from**

# There are 5 to choose from

- Bitcask

# There are 5 to choose from

- Bitcask
- Innostore

# There are 5 to choose from

- Bitcask
- Innostore
- LevelDB

# There are 5 to choose from

- Bitcask
- Innostore
- LevelDB
- Memory

# There are 5 to choose from

- Bitcask
- Innostore
- LevelDB
- Memory
- Multi



# Bitcask

# Bitcask

- A fast, append-only key-value store

# Bitcask

- A fast, append-only key-value store
- In memory key lookup table  
(key\_dir)

# Bitcask

- A fast, append-only key-value store
- In memory key lookup table (key\_dir)
- Closed files are immutable

# Bitcask

- A fast, append-only key-value store
- In memory key lookup table (key\_dir)
- Closed files are immutable
- Merging cleans up old data

# Bitcask

- A fast, append-only key-value store
- In memory key lookup table (key\_dir)
- Closed files are immutable
- Merging cleans up old data
- Apache 2 license

# Innostore

# Innostore

- Based on Embedded InnoDB



# Innostore

- Based on Embedded InnoDB
- Write ahead log plus B-tree storage

# Innostore

- Based on Embedded InnoDB
- Write ahead log plus B-tree storage
- Similar characteristics to the MySQL InnoDB plugin

# Innostore

- Based on Embedded InnoDB
- Write ahead log plus B-tree storage
- Similar characteristics to the MySQL InnoDB plugin
- Manages which pages of the B-trees are in memory

# Innostore

- Based on Embedded InnoDB
- Write ahead log plus B-tree storage
- Similar characteristics to the MySQL InnoDB plugin
- Manages which pages of the B-trees are in memory
- GPL license

# LevelDB

# LevelDB

- A Google developed key-value store

# LevelDB

- A Google developed key-value store
- Append-only

# LevelDB

- A Google developed key-value store
- Append-only
- Multiple levels of SSTable-like structures



# LevelDB

- A Google developed key-value store
- Append-only
- Multiple levels of SSTable-like structures
- BSD license

# Cluster Capacity Planning

# Cluster Capacity Planning

# Cluster Capacity Planning

- Overview

# Cluster Capacity Planning

- Overview
  - There are *MANY* variables that affect performance.

# Cluster Capacity Planning

- Overview
  - There are *MANY* variables that affect performance.
  - Safest route is to simulate load and see what happens.

# Cluster Capacity Planning

- Overview
  - There are *MANY* variables that affect performance.
  - Safest route is to simulate load and see what happens.
  - basho\_bench can help

# Cluster Capacity Planning

- Overview
  - There are *MANY* variables that affect performance.
  - Safest route is to simulate load and see what happens.
  - basho\_bench can help
- Documentation



# Cluster Capacity Planning

- Overview
  - There are MANY variables that affect performance.
  - Safest route is to simulate load and see what happens.
  - basho\_bench can help
- Documentation
  - <http://wiki.basho.com/Cluster-Capacity-Planning.html>

# Cluster Capacity Planning

- Overview
  - There are MANY variables that affect performance.
  - Safest route is to simulate load and see what happens.
  - basho\_bench can help
- Documentation
  - <http://wiki.basho.com/Cluster-Capacity-Planning.html>
  - <http://wiki.basho.com/System-Requirements.html>

# Things to Consider

## (1/3)

# Things to Consider (1/3)

- Virtualization

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.
  - Virtualize for other factors (agility or cost)

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.
  - Virtualize for other factors (agility or cost)
- Object size



# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.
  - Virtualize for other factors (agility or cost)
- Object size
  - Too small means unnecessary disk and network overhead

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.
  - Virtualize for other factors (agility or cost)
- Object size
  - Too small means unnecessary disk and network overhead
  - Too large means chunky reads and writes.

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.
  - Virtualize for other factors (agility or cost)
- Object size
  - Too small means unnecessary disk and network overhead
  - Too large means chunky reads and writes.
- Read/Write Ratio

# Things to Consider (1/3)

- Virtualization
  - Riak's workload is I/O bound
  - Run on bare metal for best performance.
  - Virtualize for other factors (agility or cost)
- Object size
  - Too small means unnecessary disk and network overhead
  - Too large means chunky reads and writes.
- Read/Write Ratio
  - Riak is not meant for dumping many tiny datapoints.

# Things to Consider (2/3)

# Things to Consider (2/3)

- “Randomness” of Reads / Writes

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
- Recently accessed objects are cached in memory.

# Things to Consider

## (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.



# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM
  - Expands the size of cached working set, lowers latency.

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM
  - Expands the size of cached working set, lowers latency.
- Disk Speed

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM
  - Expands the size of cached working set, lowers latency.
- Disk Speed
  - Faster disk means lower read/write latency.

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM
  - Expands the size of cached working set, lowers latency.
- Disk Speed
  - Faster disk means lower read/write latency.
- Number of Processors

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM
  - Expands the size of cached working set, lowers latency.
- Disk Speed
  - Faster disk means lower read/write latency.
- Number of Processors
  - More processors == more Map/Reduce throughput.

# Things to Consider (2/3)

- “Randomness” of Reads / Writes
  - Recently accessed objects are cached in memory.
  - A small working set of objects will be fast.
- Amount of RAM
  - Expands the size of cached working set, lowers latency.
- Disk Speed
  - Faster disk means lower read/write latency.
- Number of Processors
  - More processors == more Map/Reduce throughput.
  - Faster processors == lower Map/Reduce latency.

# Things to Consider

## (3/3)



# Things to Consider (3/3)

- Features

# Things to Consider (3/3)

- Features
  - RiakKV vs. Riak 2I vs. Riak Search

# Things to Consider (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends

# Things to Consider

## (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB

# Things to Consider

## (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB
  - Space/Time tradeoff

# Things to Consider

## (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB
  - Space/Time tradeoff
- Partitions

# Things to Consider (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB
  - Space/Time tradeoff
- Partitions
  - Depends on eventual number of boxes.

# Things to Consider

## (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB
  - Space/Time tradeoff
- Partitions
  - Depends on eventual number of boxes.
- Protocol



# Things to Consider

## (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB
  - Space/Time tradeoff
- Partitions
  - Depends on eventual number of boxes.
- Protocol
  - HTTP vs. Protocol Buffers

# Things to Consider

## (3/3)

- Features
  - RiakKV vs. Riak 2l vs. Riak Search
- Backends
  - Bitcask vs. Innostore vs. eLevelDB
  - Space/Time tradeoff
- Partitions
  - Depends on eventual number of boxes.
- Protocol
  - HTTP vs. Protocol Buffers
  - All clients support HTTP, some support PB.

# Performance Tuning

# Performance Tuning

- General Tips

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs
    - The More RAM the Better

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs
    - The More RAM the Better
    - Fast Disk



# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs
    - The More RAM the Better
    - Fast Disk
  - Use noatime mounts

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs
    - The More RAM the Better
    - Fast Disk
  - Use noatime mounts
  - Limit List Keys and Full Bucket MapReduce

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs
    - The More RAM the Better
    - Fast Disk
  - Use noatime mounts
  - Limit List Keys and Full Bucket MapReduce
  - Benchmark in Advance

# Performance Tuning

- General Tips
  - Start with a “DB” Like Machine Profile
    - Multi-core 64-bit CPUs
    - The More RAM the Better
    - Fast Disk
  - Use noatime mounts
  - Limit List Keys and Full Bucket MapReduce
  - Benchmark in Advance
  - Graph Everything

# Performance Tuning

# Performance Tuning

- Bitcask Backend

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)



# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)
  - Merge Trigger Settings

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)
  - Merge Trigger Settings
  - Scheduling Bitcask Merges

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)
  - Merge Trigger Settings
  - Scheduling Bitcask Merges
  - Key Expiration

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)
  - Merge Trigger Settings
  - Scheduling Bitcask Merges
  - Key Expiration
- Documentation

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)
  - Merge Trigger Settings
  - Scheduling Bitcask Merges
  - Key Expiration
- Documentation
  - <http://wiki.basho.com/Bitcask.html>

# Performance Tuning

- Bitcask Backend
  - KeyDir is in Memory
  - Filesystem Cache (Access Profile Considerations)
  - Merge Trigger Settings
  - Scheduling Bitcask Merges
  - Key Expiration
- Documentation
  - <http://wiki.basho.com/Bitcask.html>
  - <http://wiki.basho.com/Bitcask-Capacity-Planning.html>

# Performance Tuning

# Performance Tuning

- Innostore Backend



# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size
  - o\_direct

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size
  - o\_direct
  - log\_files\_in\_groups

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size
  - o\_direct
  - log\_files\_in\_groups
  - Separate Spindles for Log and Data

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size
  - o\_direct
  - log\_files\_in\_groups
  - Separate Spindles for Log and Data
- Documentation

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size
  - o\_direct
  - log\_files\_in\_groups
  - Separate Spindles for Log and Data
- Documentation
  - <http://wiki.basho.com/Setting-Up-Innostore.html>

# Performance Tuning

- Innostore Backend
  - Similar to MySQL + InnoDB
  - buffer pool size
  - o\_direct
  - log\_files\_in\_groups
  - Separate Spindles for Log and Data
- Documentation
  - <http://wiki.basho.com/Setting-Up-Innostore.html>
  - <http://wiki.basho.com/Innostore-Configuration-and-Tuning.html>



# basho\_bench

# basho\_bench

- What Does It Do?

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters
  - Test Harness / Key Generator / Value Generator

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters
  - Test Harness / Key Generator / Value Generator
  - Read / write / update mix.

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters
  - Test Harness / Key Generator / Value Generator
  - Read / write / update mix.
  - Number of workers.



# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters
  - Test Harness / Key Generator / Value Generator
  - Read / write / update mix.
  - Number of workers.
  - Worker rate. (X per second vs. as fast as possible)

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters
  - Test Harness / Key Generator / Value Generator
  - Read / write / update mix.
  - Number of workers.
  - Worker rate. (X per second vs. as fast as possible)
- Documentation

# basho\_bench

- What Does It Do?
  - Generates load against a Riak cluster. Logs and graphs results.
  - Intended for cluster capacity planning.
- Major Parameters
  - Test Harness / Key Generator / Value Generator
  - Read / write / update mix.
  - Number of workers.
  - Worker rate. (X per second vs. as fast as possible)
- Documentation
  - <http://wiki.basho.com/Benchmarking-with-Basho-Bench.html>


# basho\_bench: Configuration

```
{mode, max}.
{duration, 20}.
{concurrent, 3}.
{driver, basho_bench_driver_http_raw}.
{code_paths, ["deps/ibrowse"]}.
{key_generator, {int_to_bin,
{uniform_int, 500000000}}}.
{value_generator, {fixed_bin, 5000}}.
{http_raw_port, 8098}.
{operations, [{insert, 1}]}.
{source_dir, "foo"}.
```

# basho\_bench: Results

```
elapsed, window, total, successful, failed
10, 10, 10747, 10747, 0
20, 10, 13295, 13295, 0
30, 10, 11840, 11840, 0
40, 10, 1688, 1688, 0
50, 10, 10675, 10675, 0
60, 10, 13599, 13599, 0
70, 10, 7747, 7747, 0
80, 10, 4089, 4089, 0
90, 10, 13851, 13851, 0
100, 10, 13374, 13374, 0
110, 10, 2293, 2293, 0
120, 10, 9813, 9813, 0
130, 10, 13860, 13860, 0
140, 10, 8333, 8333, 0
150, 10, 3592, 3592, 0
160, 10, 14104, 14104, 0
170, 10, 13834, 13834, 0
...snip...
```

# basho\_bench: Results



# Monitoring

# Lager



# Lager

- Notes

# Lager

- Notes
  - Basho's replacement for Erlang logging tools

# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools

# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools
- `/var/log/riak/console.log`

# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools
- `/var/log/riak/console.log`
  - Console Output

# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools
- `/var/log/riak/console.log`
  - Console Output
- `/var/log/riak/error.log`

# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools
- `/var/log/riak/console.log`
  - Console Output
- `/var/log/riak/error.log`
  - Errors Reports

# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools
- `/var/log/riak/console.log`
  - Console Output
- `/var/log/riak/error.log`
  - Errors Reports
- `/var/log/riak/crash.log`



# Lager

- Notes
  - Basho's replacement for Erlang logging tools
  - It plays nice with your tools
- `/var/log/riak/console.log`
  - Console Output
- `/var/log/riak/error.log`
  - Errors Reports
- `/var/log/riak/crash.log`
  - Verbose Crash Output

# Logs

```
{lager, [
    %% What handlers to install with what arguments
    {handlers, [
        {lager_console_backend, info},
        {lager_file_backend, [
            {"/log/error.log", error},
            {"/log/console.log", info}
        ]}
    ]},
    %% Whether to write a crash log, and where.
    %% Commented/omitted/undefined means no crash
    {logger,
        {crash_log, "/log/crash.log"},
        %% Maximum size in bytes of events in the crash log.
        %% Default is 64kb.
        {crash_log_size, 65536},
        %% Whether to redirect error_logger messages into
        lager - defaults to true
        {error_logger_redirect, true}
    ]},
```

# Useful Commands for Monitoring

# Useful Commands for Monitoring

- `riak-admin ping`

# Useful Commands for Monitoring

- `riak-admin ping`
- Check that a node is running.

# Useful Commands for Monitoring

- `riak-admin ping`
- Check that a node is running.
  - Also available as `localhost:8098/ping`

# Useful Commands for Monitoring

- `riak-admin ping`
- Check that a node is running.
  - Also available as `localhost:8098/ping`
- `riak-admin status`

# Useful Commands for Monitoring

- `riak-admin ping`
- Check that a node is running.
  - Also available as `localhost:8098/ping`
- `riak-admin status`
  - Show status of the node.



# Useful Commands for Monitoring

- `riak-admin ping`
- Check that a node is running.
  - Also available as `localhost:8098/ping`
- `riak-admin status`
  - Show status of the node.
  - Also available as `localhost:8098/stats`

# riak-admin status

```
1-minute stats for 'riak@10.0.201.105'
```

```
-----  
vnode gets : 3055  
vnode_puts : 1525  
read_repairs : 1  
vnode_gets_total : 14211579  
vnode_puts_total : 7158758  
node_gets : 1015  
node_gets_total : 4761211  
node_get_fsm_time_mean : 103288  
node_get_fsm_time_median : 1143  
node_get_fsm_time_95 : 711812  
node_get_fsm_time_99 : 1212212  
node_get_fsm_time_100 : 2574281  
node_puts : 507  
node_puts_total : 2378135  
node_put_fsm_time_mean : 22447  
node_put_fsm_time_median : 1226  
node_put_fsm_time_95 : 162412  
node_put_fsm_time_99 : 493482  
node_put_fsm_time_100 : 833879  
read_repairs_total : 26246
```

# Other Tools

# Other Tools

- Unix Commands

# Other Tools

- Unix Commands
  - top

# Other Tools

- Unix Commands
  - top
  - iostat

# Other Tools

- Unix Commands
  - top
  - iostat
- Enterprise Tools

# Other Tools

- Unix Commands
  - top
  - iostat
- Enterprise Tools
  - JMX Monitoring



# Other Tools

- Unix Commands
  - top
  - iostat
- Enterprise Tools
  - JMX Monitoring
  - SNMP Monitoring

# Backups

# Backups

- Overview

# Backups

- Overview
  - Backup options depend on backend.

# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index

# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index
  - Immutable files.

# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index
  - Immutable files.
  - Can use rsync, snapshots, cp, or riak-admin backup.

# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index
  - Immutable files.
  - Can use rsync, snapshots, cp, or riak-admin backup.
- Innostore



# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index
  - Immutable files.
  - Can use rsync, snapshots, cp, or riak-admin backup.
- Innostore
  - Only riak-admin backup

# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index
  - Immutable files.
  - Can use rsync, snapshots, cp, or riak-admin backup.
- Innostore
  - Only riak-admin backup
- Enterprise Replication

# Backups

- Overview
  - Backup options depend on backend.
- Bitcask, LevelDB, & Merge\_Index
  - Immutable files.
  - Can use rsync, snapshots, cp, or riak-admin backup.
- Innostore
  - Only riak-admin backup
- Enterprise Replication
  - Mirror to another cluster.

# Rolling Upgrades

# Rolling Upgrades

- Overview

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps



# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.
  2. Back up configuration and data directories. (If possible.)

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.
  2. Back up configuration and data directories. (If possible.)
  3. Install new package.

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.
  2. Back up configuration and data directories. (If possible.)
  3. Install new package.
  4. Start node.

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.
  2. Back up configuration and data directories. (If possible.)
  3. Install new package.
  4. Start node.
  5. Do next node.

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.
  2. Back up configuration and data directories. (If possible.)
  3. Install new package.
  4. Start node.
  5. Do next node.
- Documentation

# Rolling Upgrades

- Overview
  - You can upgrade Riak without taking down cluster.
  - Instead, take down one node at a time. Stay read/write available.
- Steps
  1. Stop node.
  2. Back up configuration and data directories. (If possible.)
  3. Install new package.
  4. Start node.
  5. Do next node.
- Documentation
  - <http://wiki.basho.com/Rolling-Upgrades.html>

# EC2 and EBS



# EC2 and EBS

- Follow DB Best Practices

# EC2 and EBS

- Follow DB Best Practices
- To EBS or Not to EBS?

# EC2 and EBS

- Follow DB Best Practices
- To EBS or Not to EBS?
- Raid0 is Faster, but an OPS Headache

# EC2 and EBS

- Follow DB Best Practices
- To EBS or Not to EBS?
- Raid0 is Faster, but an OPS Headache
- Speed vs Fault Tolerance (Availability Zones)

# EC2 and EBS

- Follow DB Best Practices
- To EBS or Not to EBS?
- Raid0 is Faster, but an OPS Headache
- Speed vs Fault Tolerance (Availability Zones)
- Specific Tweaks for EC2 (`net_ticktime`)

# EC2 and EBS

- Follow DB Best Practices
- To EBS or Not to EBS?
- Raid0 is Faster, but an OPS Headache
- Speed vs Fault Tolerance (Availability Zones)
- Specific Tweaks for EC2 (net\_ticktime)
- There are MANY successful deployments

# Troubleshooting

# Node Crashed, Now What? (1/3)



# Node Crashed, Now What? (1/3)

- Overview

# Node Crashed, Now What? (1/3)

- Overview
  - Approach this like any other troubleshooting.

# Node Crashed, Now What? (1/3)

- Overview
  - Approach this like any other troubleshooting.
  - Stay calm. Your system is still fully available.

# Node Crashed, Now What? (1/3)

- Overview
  - Approach this like any other troubleshooting.
  - Stay calm. Your system is still fully available.
  - Survey the scene:

# Node Crashed, Now What? (1/3)

- Overview
  - Approach this like any other troubleshooting.
  - Stay calm. Your system is still fully available.
  - Survey the scene:
    - Was anything abnormal? High load? Notice any patterns?

# Node Crashed, Now What? (2/3)

# Node Crashed, Now What? (2/3)

- Save Log Files!

# Node Crashed, Now What? (2/3)

- Save Log Files!
  - Ensure you can find the root cause later.



# Node Crashed, Now What? (2/3)

- Save Log Files!
  - Ensure you can find the root cause later.
  - `/var/log/riak`

# Node Crashed, Now What? (2/3)

- Save Log Files!
  - Ensure you can find the root cause later.
  - `/var/log/riak`
  - `/var/lib/riak/ring`

# Node Crashed, Now What? (3/3)

# Node Crashed, Now What? (3/3)

- Did you run out of resources?

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?
  - Fix hardware, restore from backup (if available/necessary)

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?
  - Fix hardware, restore from backup (if available/necessary)
  - Remember to re-ip!



# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?
  - Fix hardware, restore from backup (if available/necessary)
  - Remember to re-ip!
- Software Error / Some Other Reason?

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?
  - Fix hardware, restore from backup (if available/necessary)
  - Remember to re-ip!
- Software Error / Some Other Reason?
  - Contact mailing list or get support.

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?
  - Fix hardware, restore from backup (if available/necessary)
  - Remember to re-ip!
- Software Error / Some Other Reason?
  - Contact mailing list or get support.
- Documentation

# Node Crashed, Now What? (3/3)

- Did you run out of resources?
  - Plan to scale out.
- Did your hardware fail?
  - Fix hardware, restore from backup (if available/necessary)
  - Remember to re-ip!
- Software Error / Some Other Reason?
  - Contact mailing list or get support.
- Documentation
  - <http://wiki.basho.com/Recovering-a-failed-node.html>