

# Data Modeling for Scale with Riak Data Types

Sean Cribbs

@seancribbs #riak #datatypes

GlueCon 2014


I work for Basho  
We make  riak

Visit our  
booth



**Riak is  
Eventually Consistent**

# Eventual Consistency




Replicated  
Loose coordination  
Convergence

# Eventual is Good


- ✓ Fault-tolerant
- ✓ Highly available
- ✓ Low-latency

# Consistency?



No clear winner!  
Throw one out?  
Keep both?

# Consistency?




No clear winner!

Throw one out?

Keep both?

Cassandra

# Consistency?



No clear winner!

Throw one out?

Keep both?

Cassandra

Riak



# Conflicts!

A!



B!



# Semantic Resolution

- Your app knows the domain - use business rules to resolve
- Amazon Dynamo's shopping cart

# Semantic Resolution

- Your app knows the domain - use business rules to resolve
- Amazon Dynamo's shopping cart

“Ad hoc approaches have proven brittle and error-prone”


# Convergent Replicated Data Types

# Convergent Replicated Data Types

useful abstractions



# Convergent Replicated Data Types



multiple independent  
copies



useful abstractions

resolves automatically  
toward a single value

# Convergent Replicated Data Types

multiple independent  
copies

useful abstractions

# How CRDTs Work

- A partially-ordered **set of values**
- A **merge** function
- An **identity** value
- **Inflation** operations



# How CRDTs Work


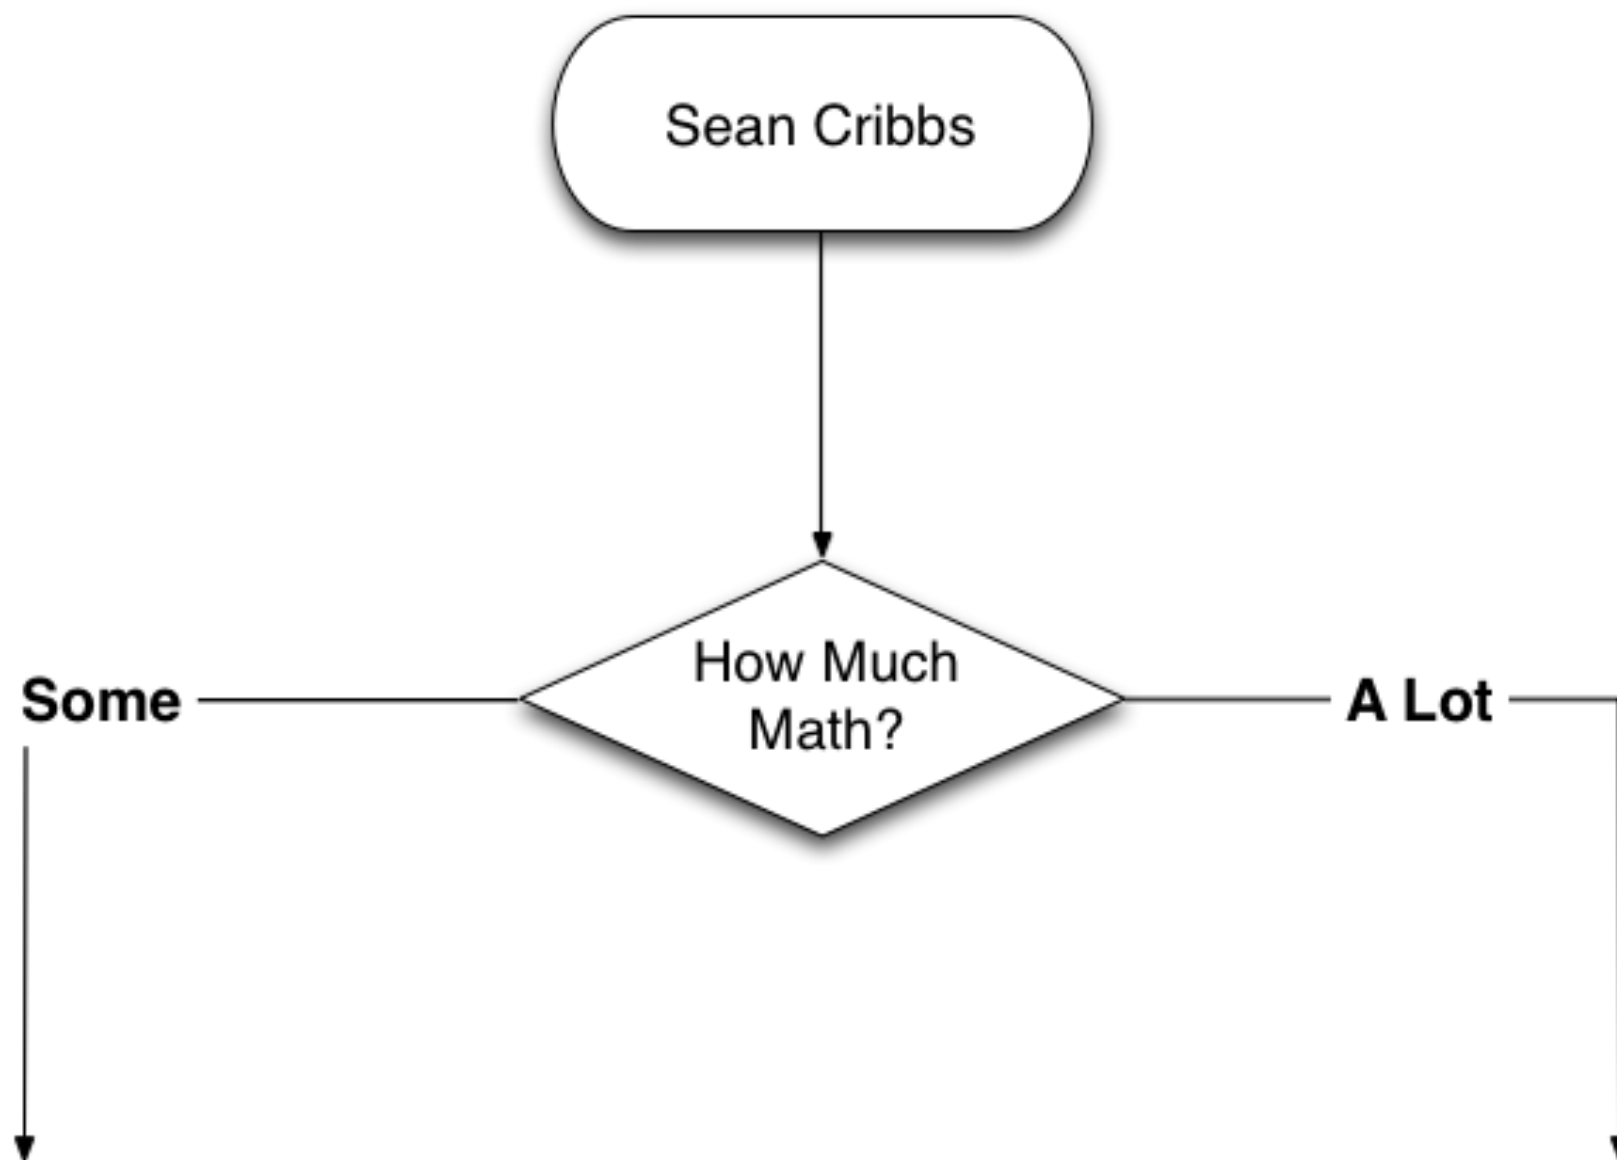
- A partially-ordered **set of values**
- A **merge** function
- An **identity** value
- **Inflation** operations

## What CRDTs Enable

- Consistency **without coordination**
- **Fluent**, rich interaction with data



This research is supported in part by European FP7 project 609 551  
SyncFree <http://syncfree.lip6.fr/> (2013--2016).



**Forget CRDTs**  
**Do Data Modeling**

# Data Modeling for Riak

- Identify needs for both **read** and **write**
- Design around **key as index**
- **Denormalize** relationships if possible
- Weigh **data size** against **coherence**

# Riak Data Types

# Riak Data Types

Counter :: int

increment

decrement

# Riak Data Types

Counter :: int

increment

decrement

Set :: { bytes }

add\*

remove



# Riak Data Types

Map :: bytes → DT

update\*

remove

Counter :: int

increment

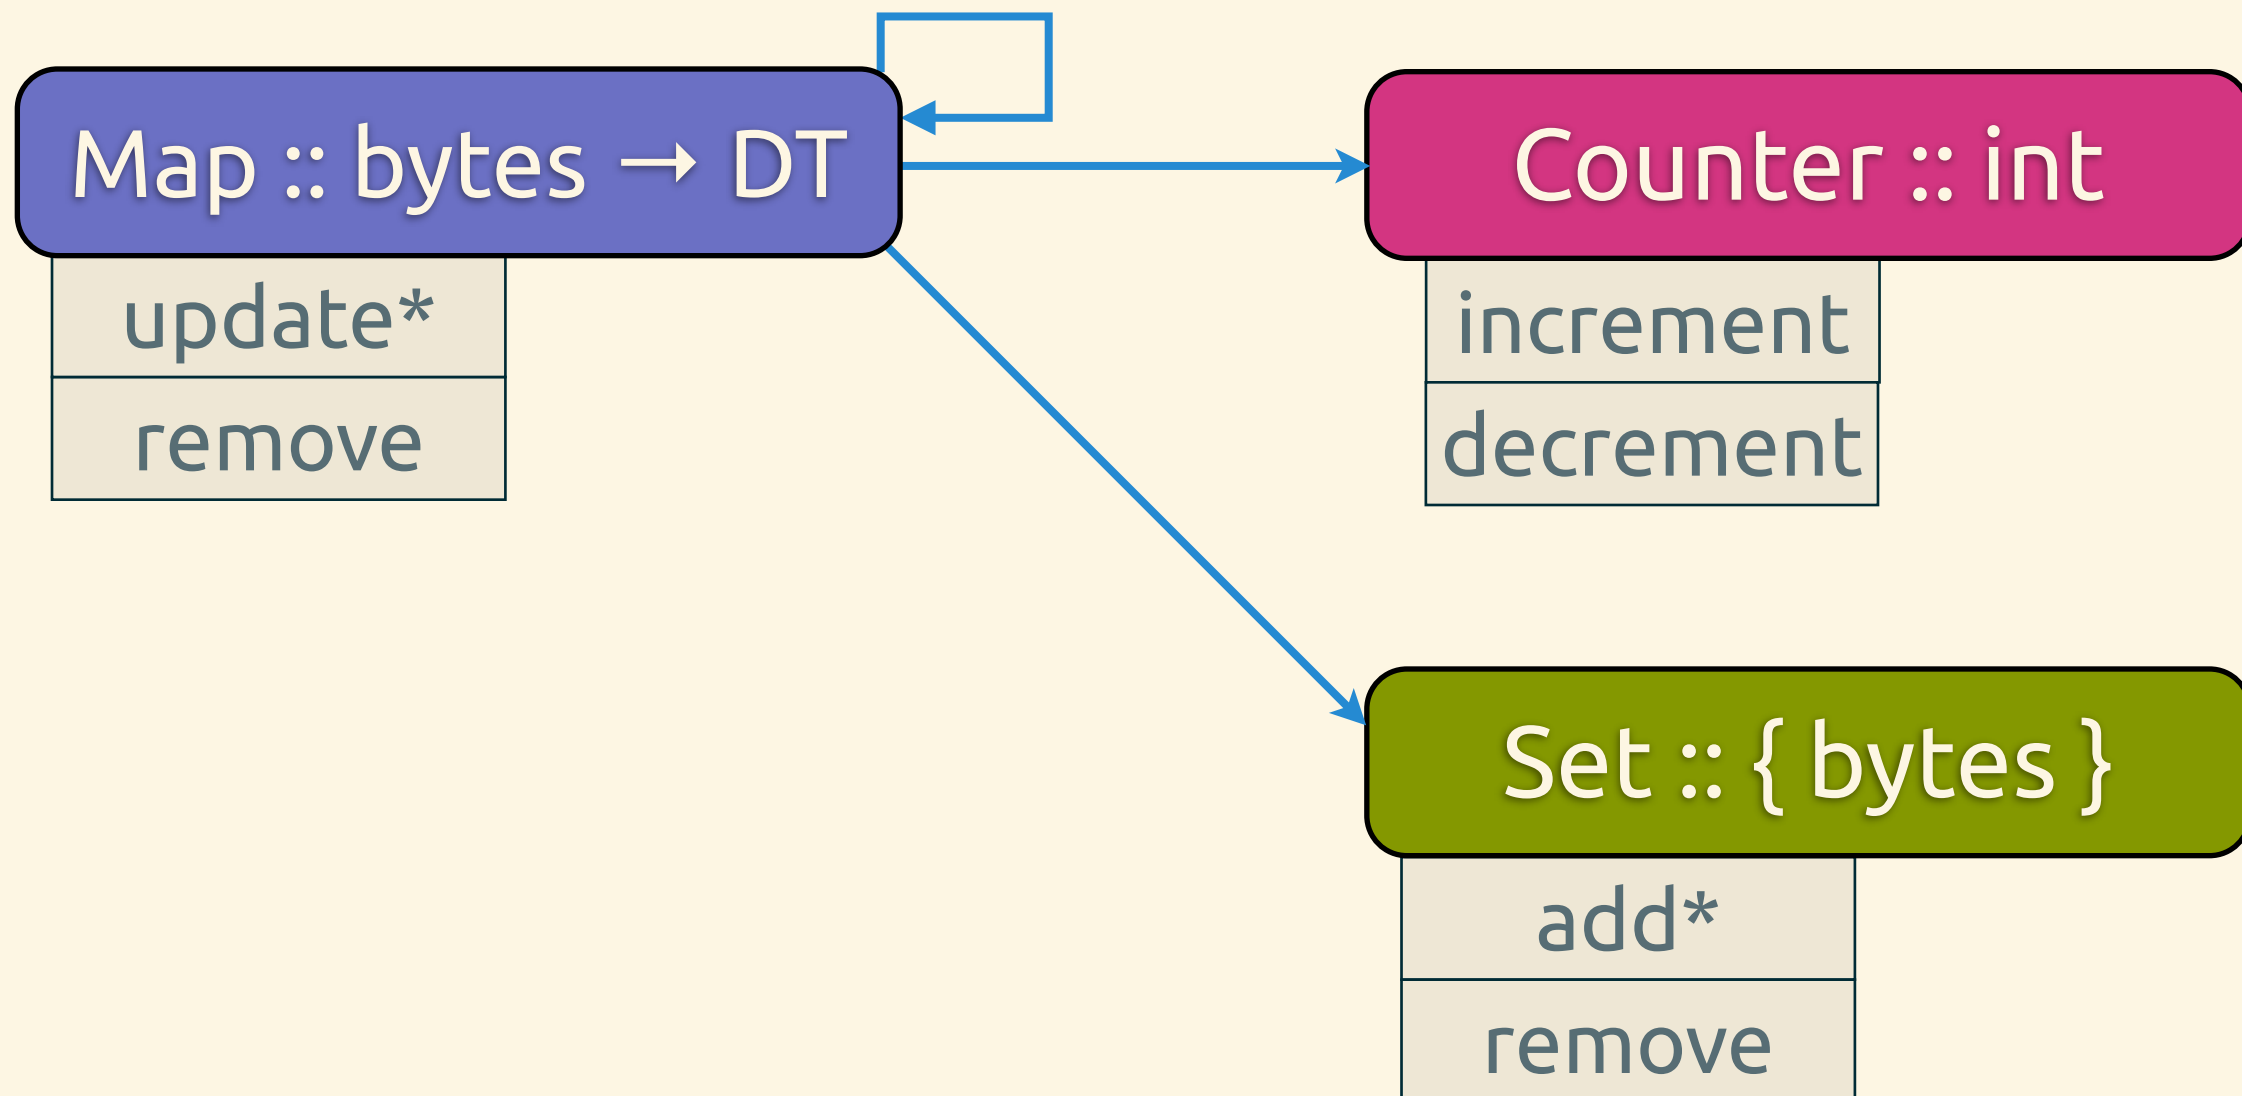
decrement

Set :: { bytes }

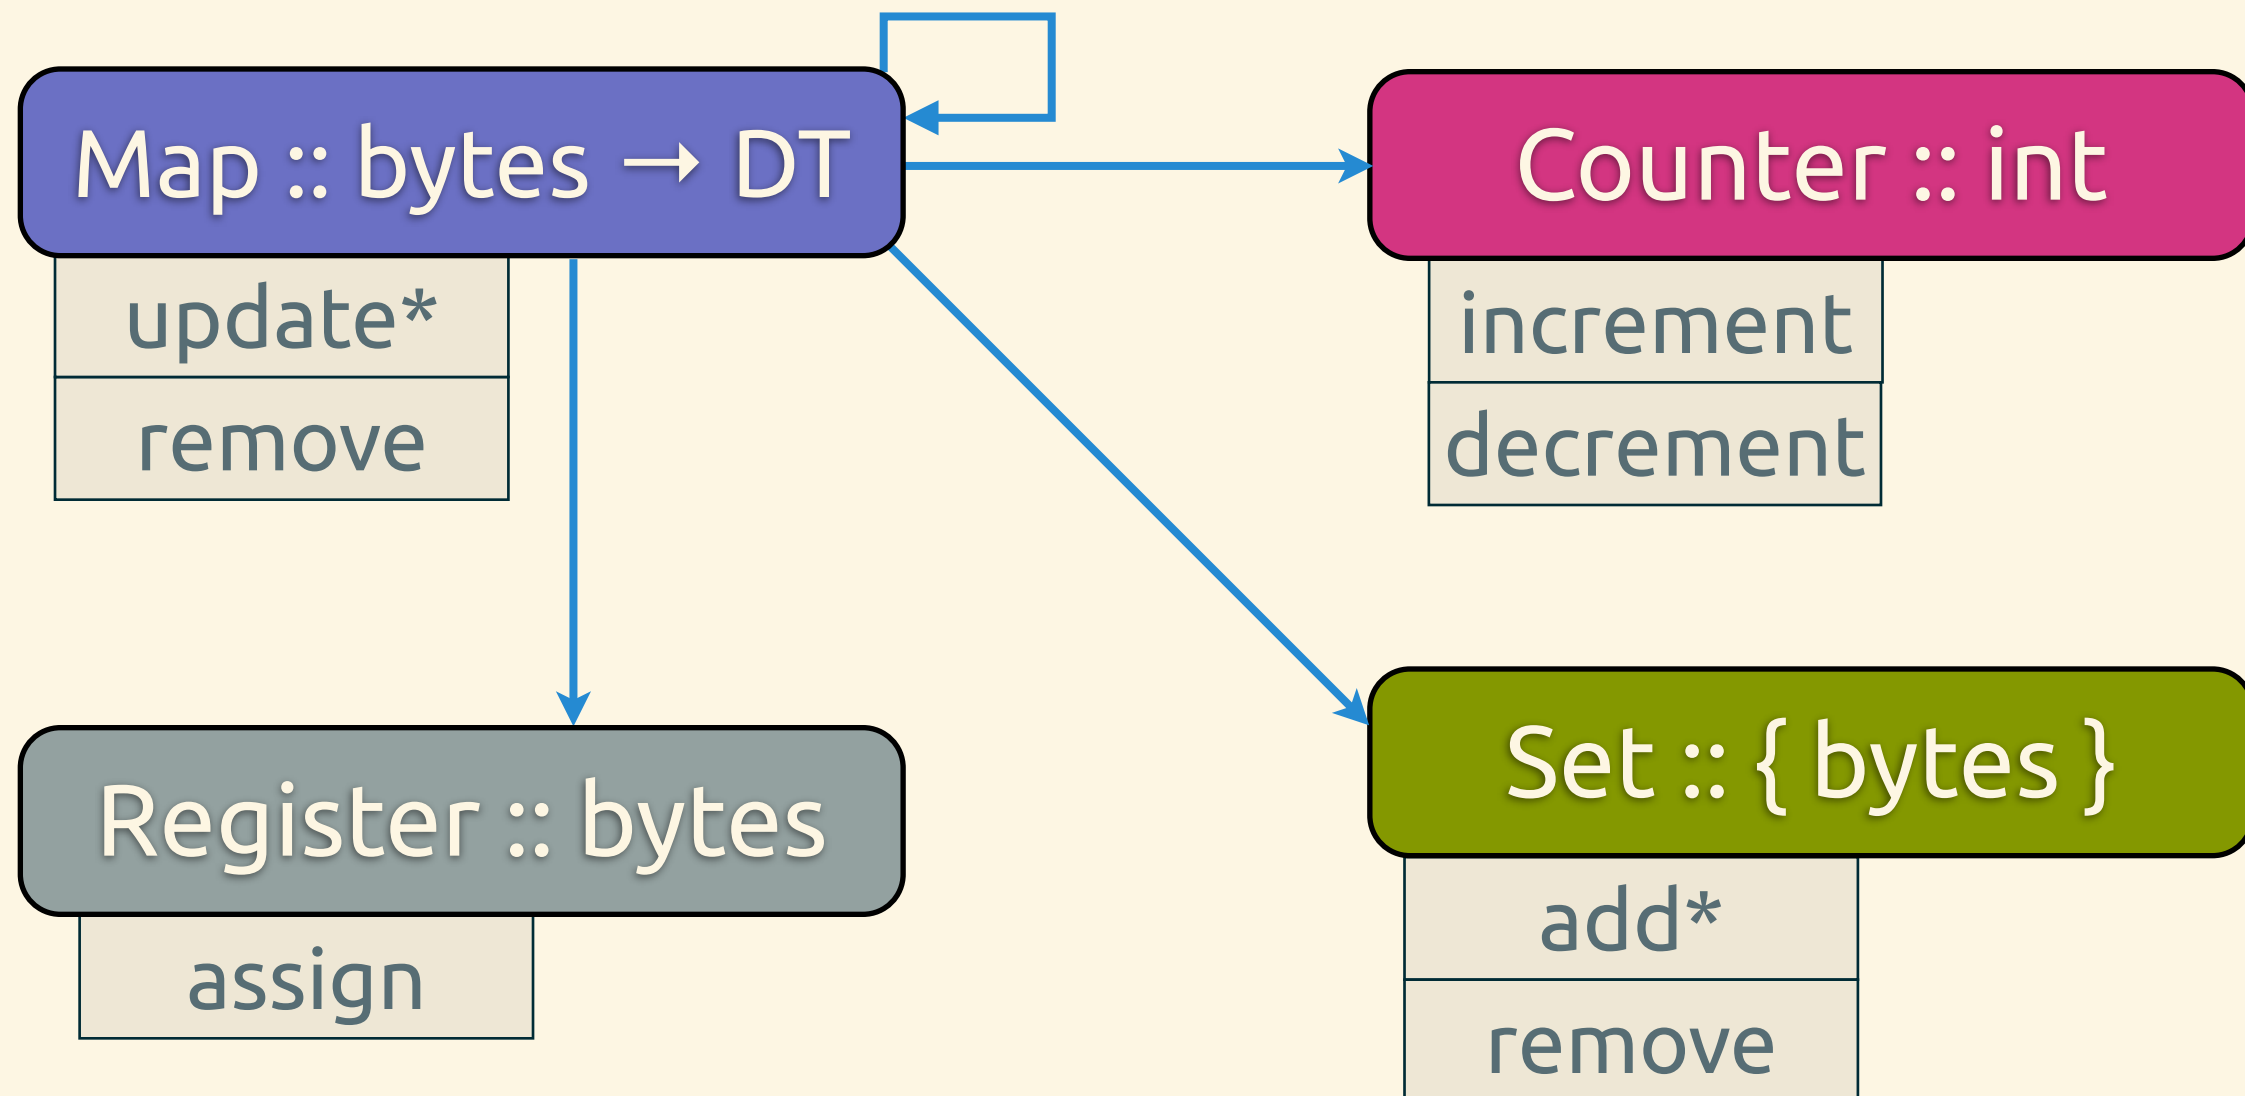
add\*

remove

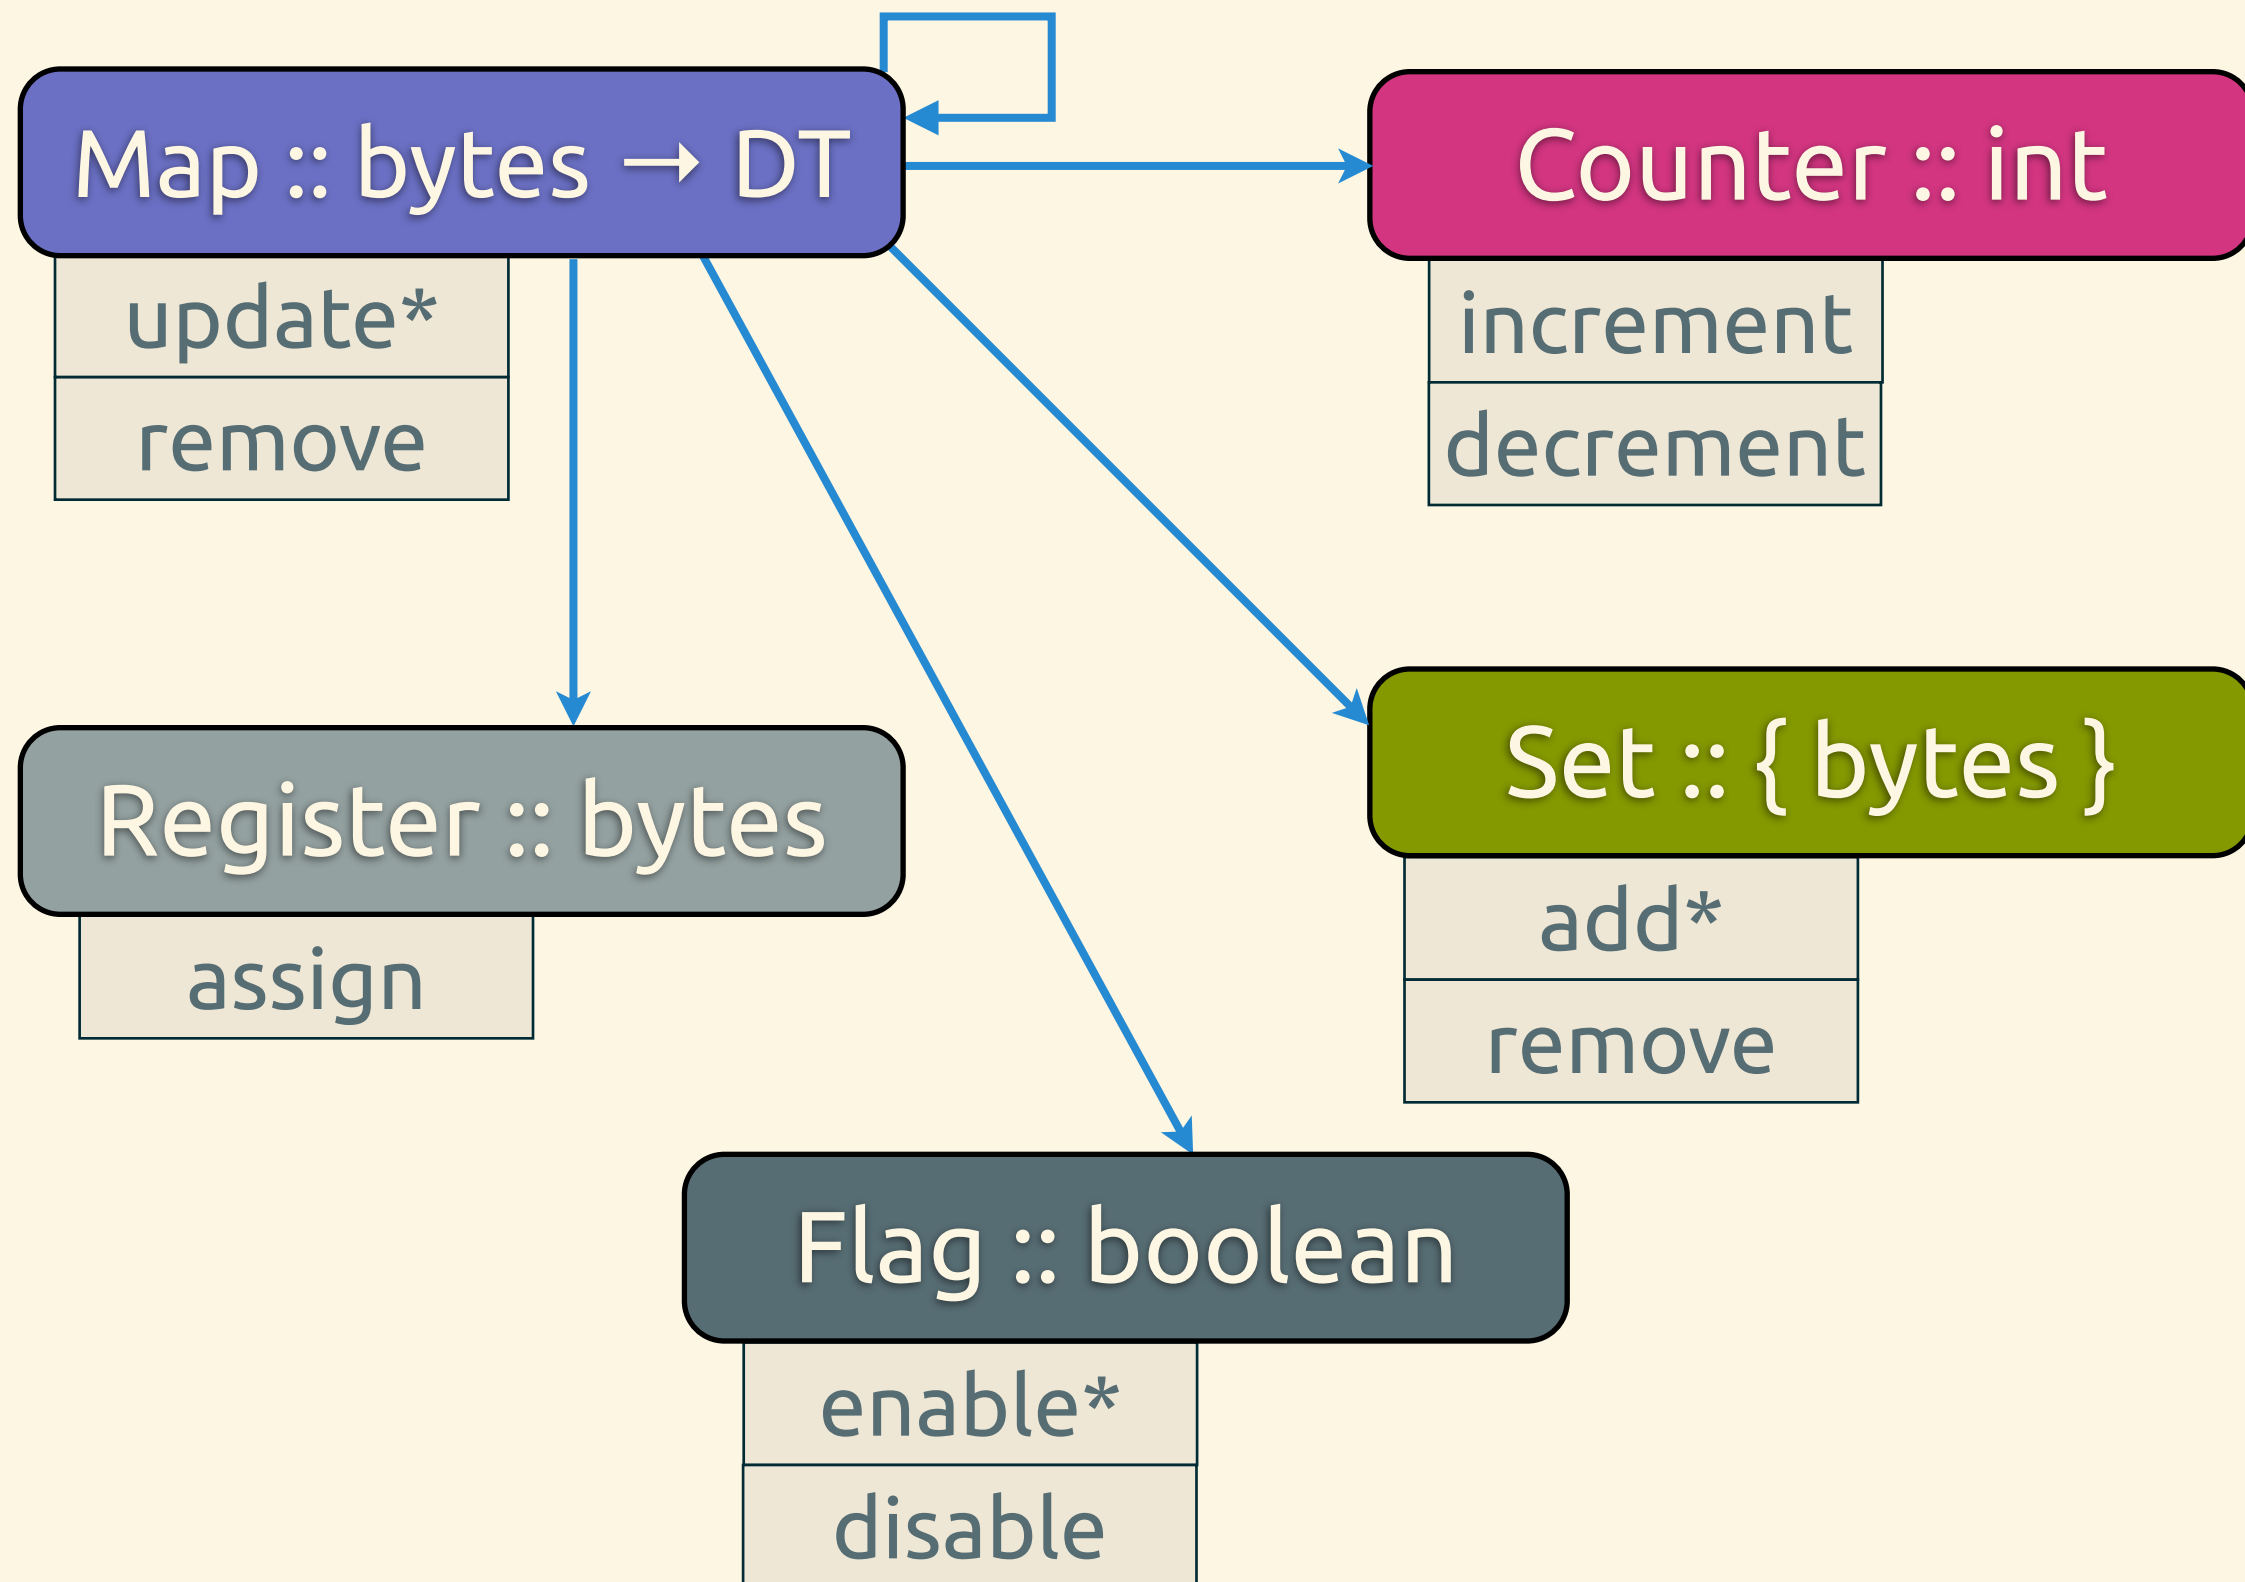
# Riak Data Types



# Riak Data Types




# Riak Data Types





**MAD DATA**





# Counters



# Ad Network

- **Impressions** - when someone sees an ad
- **Click-through** - when someone clicks on an ad
- Hourly rollups  
`ad-metrics/<campaign>/<type>-<hour>`



# Ad Network

```
$ riak-admin bucket-type create ad-metrics \  
  '{"props":{"datatype":"counter"}}'  
ad-metrics created
```

```
$ riak-admin bucket-type activate ad-metrics  
ad-metrics has been activated
```

```
$ riak-admin bucket-type list  
ad-metrics (active)
```





# Ad Network

```
from riak import RiakClient
from rogersads import RIAK_CONFIG
from time import strftime


client = RiakClient(**RIAK_CONFIG)

metrics = client.bucket_type('ad-metrics')

def record_metric(campaign, metric_type):
    key = metric_type + strftime('-%Y%m%d-%H')
    counter = metrics.bucket(campaign).new(key)
    counter.increment()
    counter.store()
```



# Ad Network





# Sets



# PartyOn

- **RSVPs** - guest lists
- **Connections** - friends lists per-user
- **Likes** - expressing interest



# PartyOn

```
$ riak-admin bucket-type create partyon-sets \  
  '{"props":{"datatype":"set"}}'  
partyon-sets created
```

```
$ riak-admin bucket-type activate partyon-sets  
partyon-sets has been activated
```

```
$ riak-admin bucket-type list  
partyon-sets (active)
```



# PartyOn

- **RSVPs**  
`partyon-sets/rsvps/<eventid>`
- **Connections**  
`partyon-sets/friends/<userid>`
- **Likes**  
`partyon-sets/likes/<eventid>`



# PartyOn

```
from riak.datatypes import Set  
  
sets = client.bucket_type('partyon-sets')  
  
rsvps = sets.bucket('rsvps')  
friends = sets.bucket('friends')  
likes = sets.bucket('likes')
```



# PartyOn

```
def rsvp_get(event):  
    return rsvps.get(event) # Returns a Set  
  
def rsvp_add(event, user):  
    guests = rsvps.new(event)  
    guests.add(user)  
    guests.store(return_body=True)  
    return guests.context  
  
def rsvp_remove(event, user, context):  
    guests = Set(rsvps, event, context=context)  
    guests.remove(user)  
    guests.store()
```





# Maps

(and the rest)



# GameNet

- **User profiles** - demographic data  
`users/<userid>/profile`
- **Achievements** - trophies per game  
`users/<userid>/trophies`
- **Game state** - progress and stats  
`users/<userid>/<gameid>`



# GameNet

```
$ riak-admin bucket-type create users \  
  '{"props":{"datatype":"map"}}'  
users created
```

```
$ riak-admin bucket-type activate users  
users has been activated
```

```
$ riak-admin bucket-type list  
users (active)
```



# GameNet

```
users = client.bucket_type('users')
```

```
def update_profile(user, fields):  
    profile = users.bucket(user).get('profile')  
  
    for field in fields:  
        if field in USER_FLAGS:  
            if fields[field]:  
                profile.flags[field].enable()  
            else:  
                profile.flags[field].disable()  
        else:  
            value = fields[field]  
            profile.registers[field].assign(value)  
  
    profile.store()
```





# GameNet

```
def add_trophy(user, game, trophy):  
    trophies = users.bucket(user).get('trophies')  
  
    trophies.sets[game].add(trophy)  
    trophies.store()  
  
def get_trophies(user, game):  
    trophies = users.bucket(user).get('trophies')  
    return trophies.sets[game].value
```



# GameNet

```
def build_structure(user, game, structure, gold,
                    wood, stone):
    gamestate = users.bucket(user).get(game)
    gamestate.sets['structures'].add(structure)
    gamestate.counters['gold'].decrement(gold)
    gamestate.counters['wood'].decrement(wood)
    gamestate.counters['stone'].decrement(stone)
    gamestate.store(return_body=True)
    return gamestate.value
```

# Benefits

- **Richer interactions**, familiar types
- Write **mutations**, not state
- No **merge function** to write
- Same **reliability** and **predictability** of vanilla Riak

# Caveats

- **Value size** still matters
- Updates **not idempotent**
- Cross-key atomicity **not possible** (yet)



# Future

- Riak 2.0 due out this summer - betas available now!
- Richer querying, lighter storage requirements, more types

# C.R.D.T.

TO RESIST BUGS  
AND WRITE LOSS.

