

Riak Tools & Cluster Operations

Tools: riak-admin

test

```
## performs a read/write test against the cluster  
riak@node:~$ riak-admin test  
Successfully completed 1 read/write cycle to  
'riak@node.example.com'
```



js-reload

- Forces the embedded Javascript virtual machines to be restarted

wait-for-service

```
## useful for scripting to check when backend is ready  
riak@node:~$ riak-admin wait-for-service riak_kv riak@<nodename>
```



ringready

```
## checks if all nodes in the cluster agree on the ring state  
riak@node:~$ riak-admin ringready
```



transfers

```
## checks for ownership or partition handoff  
riak@node:~$ riak-admin transfers
```

```
## example output
```

```
Active Transfers:
```

```
'dev5@127.0.0.1' waiting to handoff 39 partitions  
'dev4@127.0.0.1' waiting to handoff 39 partitions  
'dev3@127.0.0.1' waiting to handoff 38 partitions
```



transfer-limit

- Change the handoff_concurrency limit
- Limited from 2-8
- Affects node performance of other ops
- Remember to change it back

cluster-info

```
## output information from all nodes to /tmp/cluster_info.txt
```

```
riak@node:~$ riak-admin cluster_info /tmp/cluster_info.txt
```

```
## output information from the current node
```

```
riak@node:~$ riak-admin cluster_info /tmp/cluster_info.txt local
```

```
## output information from a subset of nodes
```

```
riak@node:~$ riak-admin cluster_info /tmp/cluster_info.txt
```

```
riak@192.168.1.10 riak@192.168.1.11
```



member-status

- Prints the current status of all cluster members

ring-status

- Ring ready
- Ownership handoff
- Unreachable nodes

vnode-status

- Per vnode information
- Backend file information

aae-status

- Exchanges
- Entropy Trees
- Keys Repaired

diag

```
## checks riak configuration and performs some system diagnosis  
riak@node:~$ riak-admin diag
```

```
## example output
```

```
[critical] vm.swappiness is 60, should be no more than 0  
[critical] net.core.wmem_default is 229376, should be at least  
8388608  
[critical] net.core.rmem_default is 229376, should be at least  
8388608  
[critical] net.core.wmem_max is 131071, should be at least 8388608  
[notice] Data directory /var/lib/riak/bitcask is not mounted with  
'noatime'. Please remount its disk with the 'noatime' flag to  
improve performance.
```

status

- One-minute stats
 - GETs & PUTs: mean, median, 95/99/100%
 - FSM times
 - GET FSM siblings
- Totals
- Erlang CPU & memory statistics
- Miscellaneous information & versions

top

```
## check cpu, memory usage and message queues in the erlang vm  
riak@node:~$ riak-admin top [-interval N] [-sort reductions|  
memory|msg_q] [-lines N]
```



Basic Commands

```
riak-admin { test | ringready | member-status | ring-status |  
status }
```

```
riak-admin cluster { join | leave | replace | plan | commit }
```



Tools: riak-debug

Cluster Operations

Riak Cluster Operations

<http://goo.gl/89gx7e>

Rolling Restarts

```
riak-admin ringready # should report true
riak-admin transfers # should NOT show pending transfers
riak stop
riak start
riak-admin wait-for-service riak_kv
## move to next node
```

Upgrading Riak

```
riak-admin ringready # should report true
riak-admin transfers # should NOT show pending transfers
riak stop
## install new package
## make any app.config changes
riak start
riak-admin wait-for-service riak_kv
## move to next node
```



Transfers & Transfer-Limit

- When a node is unavailable, fallbacks are used
- When a node re-joins the ring hinted-handoff transfers occur
- Transfer-limit governs how quickly this happens

Adding Nodes

```
## configure the new node
riak start
riak-admin cluster join riak@<existing_node>
riak-admin cluster plan
riak-admin cluster commit
## wait for ownership handoff to complete
watch -n 1 "riak-admin transfers"
```



Removing Nodes

```
## connect to the node being removed
riak-admin cluster leave
## or request another node leave
riak-admin cluster leave riak@node6
## plan and commit
riak-admin cluster plan
riak-admin cluster commit
## wait for ownership handoff to complete
watch -n 1 "riak-admin transfers"
```

Replacing Nodes

- Backup data directory
- Install and start Riak on new node
- Plan the join of the new node to the cluster
- Plan the replacement of the existing node with the new node
- Plan and commit changes

Renaming Nodes

- Changing a nodes name or IP address can be done without cluster impact
- Down the first node (node1 a)
- Mark down from a different node (node2)
- Edit vm.args and app.config (node1 a)
- Rename the ring directory (node1 a)
- Start the node and join to cluster (node1 b)
- Force replace node1 a with node1 b

Force-Replace a Node

```
## if a node is completely unresponsive e.g hardware failure  
riak-admin cluster force-replace riak@<node4> riak@<node7>
```

