

Performance Tuning

System Configuration

- OS level tweaks
 - Adjust ulimit
 - Mount data disk with noatime
 - Swappiness can affect IO
 - Disk I/O Scheduler

Riak Tuning

```
## in /etc/riak/vm.args
```

```
+S 8:8 # half number of cores when using LevelDB
```

```
+zdbbl 32768 # if busy_dist_port_seen in logs
```



Riak Tuning

```
## in /etc/riak/app.config
```

```
{leveldb, [  
    {cache_size, 8388608},  
    {data_root, "/var/lib/riak/leveldb"},  
    {max_open_files, 100},  
    {write_buffer_size_max, 62914560},  
    {write_buffer_size_min, 31457280}  
]},
```



OS Tuning

```
## in /etc/sysctl.d/riak.conf
vm.swappiness = 0
net.ipv4.tcp_max_syn_backlog = 40000
net.core.somaxconn=4000
net.ipv4.tcp_timestamps = 0
net.ipv4.tcp_sack = 1
net.ipv4.tcp_window_scaling = 1
net.ipv4.tcp_fin_timeout = 15
net.ipv4.tcp_keepalive_intvl = 30
net.ipv4.tcp_tw_reuse = 1
## for 10Gb Networks
net.core.rmem_default = 8388608
net.core.rmem_max = 8388608
net.core.wmem_default = 8388608
net.core.wmem_max = 8388608
net.core.netdev_max_backlog = 10000
```



OS Tuning

- EXT4
noatime,nodiratime,barrier=0,data=writeback
- XFS
noatime,nodiratime,nobarrier,logbufs=8,logbsize=256k,allocsize=2M

OS Tuning

- `/sys/block/X/queue/scheduler`
 - `noop`
 - `deadline`
- `/sys/block/x/queue/nr_requests`
 - `1024`

Next steps

- docs.basho.com
- OS & filesystem settings
- app.config & vm.args
- Backend configuration and settings